



Aalto University  
School of Science

# Robustness, Reliability, Resilience and Elasticity (R3E) for Big Data/Machine Learning Systems


*Hong-Linh Truong*

*Department of Computer Science*

*[linh.truong@aalto.fi](mailto:linh.truong@aalto.fi), <https://rdsea.github.io>*

# Our focus in this course

The focus



**Any idea,  
thought,  
expectation?**

# Learning objectives

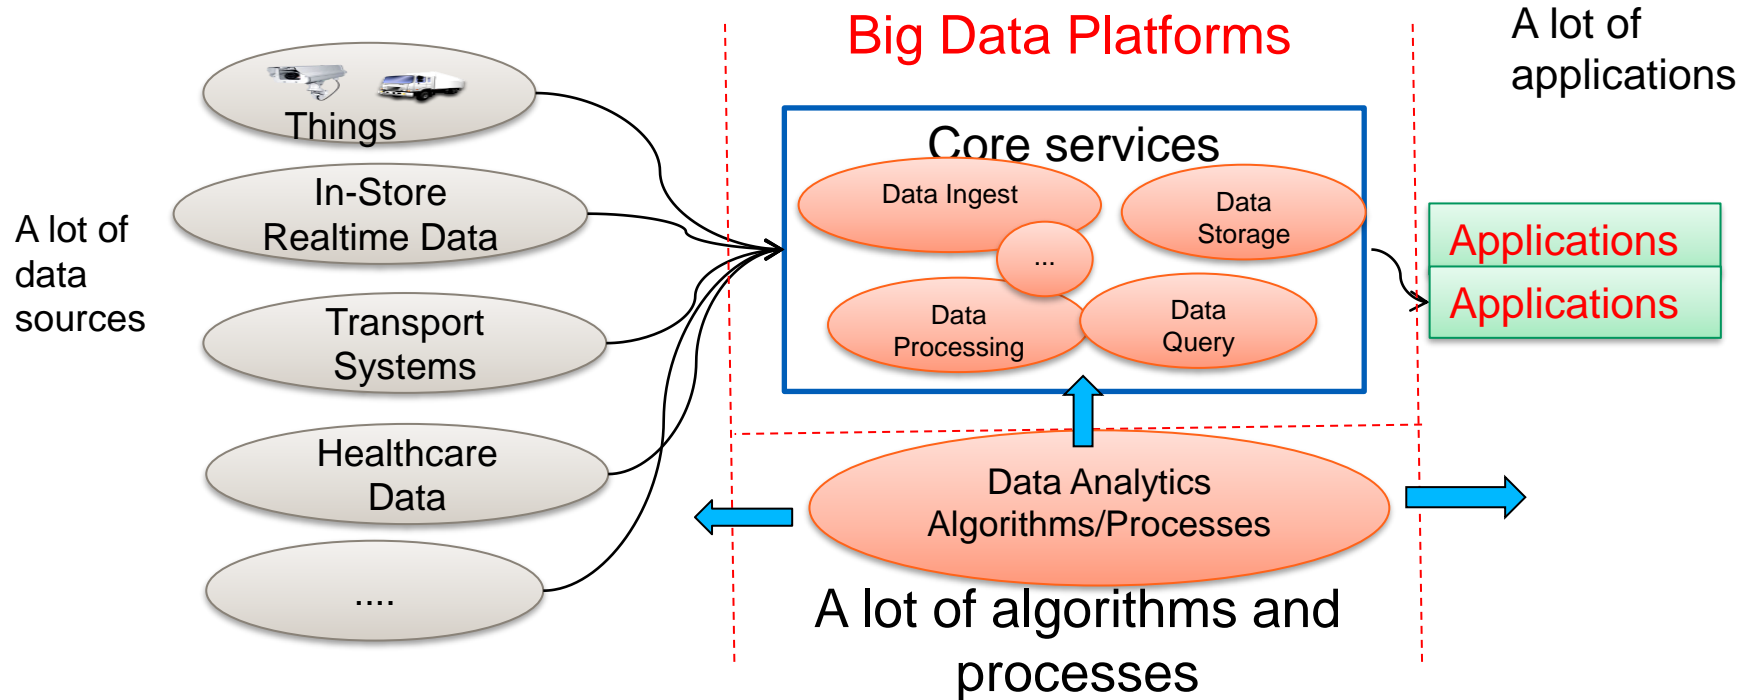
- **Identify commonality and complexity in end-to-end Big Data/ML systems**
- **Understand design goals and concerns for robustness, reliability, resilience and elasticity of Big Data/ML systems**
- **Learn an elasticity-based approach for R3E**

# Commonality and complexity in Big Data and Machine Learning systems

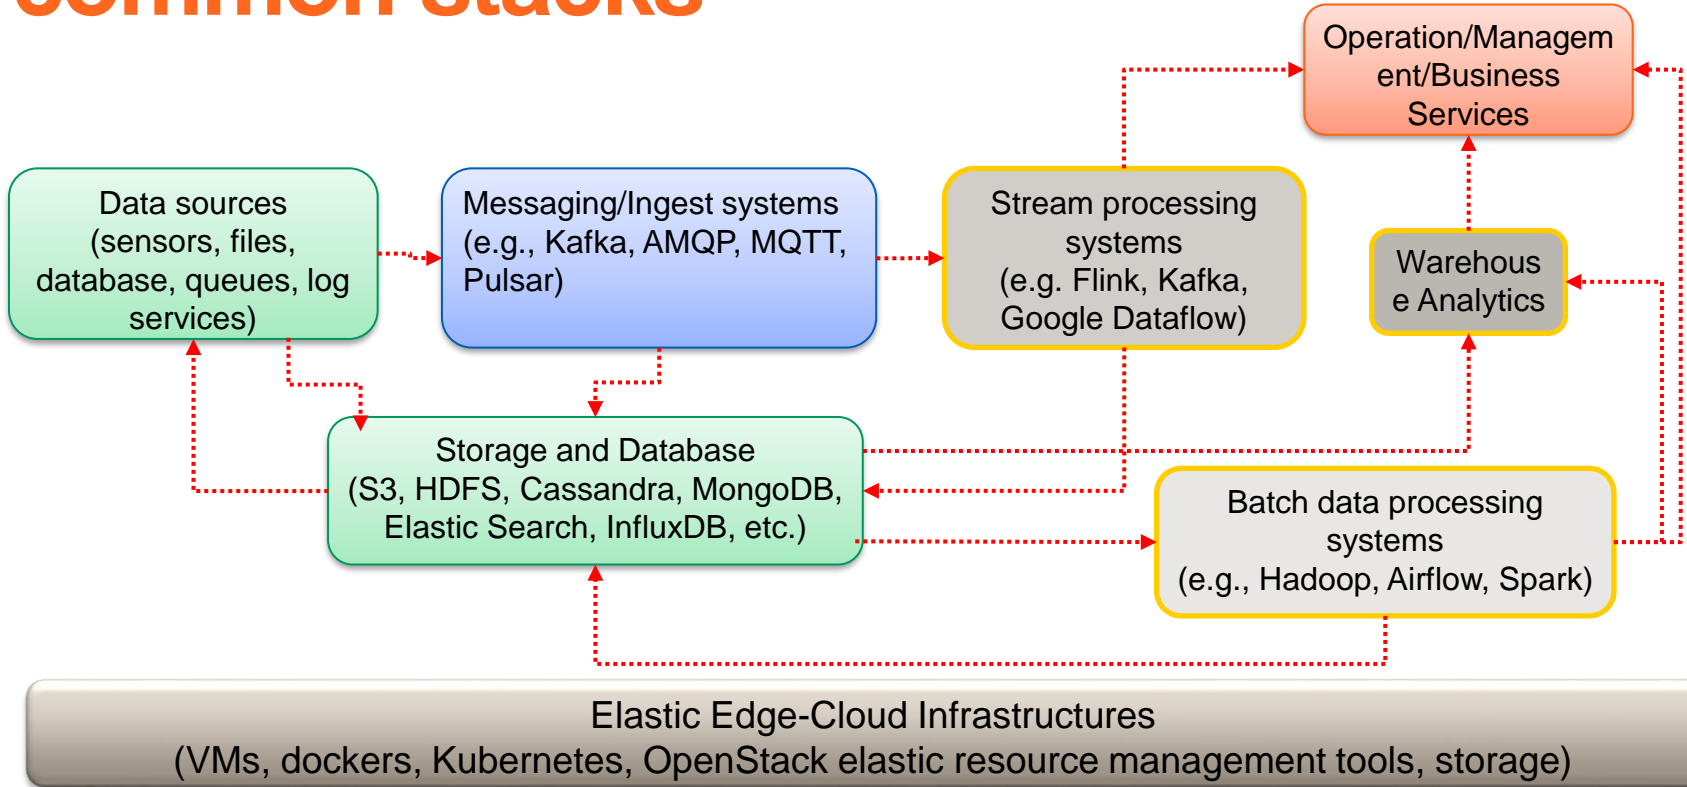
# Big data with V\*

- **Volume:**
  - big size, large data set, massive of small data
- **Variety:**
  - complexity of different formats and types of data
- **Velocity:**
  - generating speed, data movement speed
- **Veracity:**
  - quality is very different (timeliness, accuracy, etc.)

# A bird view of big data platforms



# Big data at large-scale: example of common stacks



# Examples from Big Data Platforms

**<https://version.aalto.fi/gitlab/bigdataplatforms/cs-e4640>**

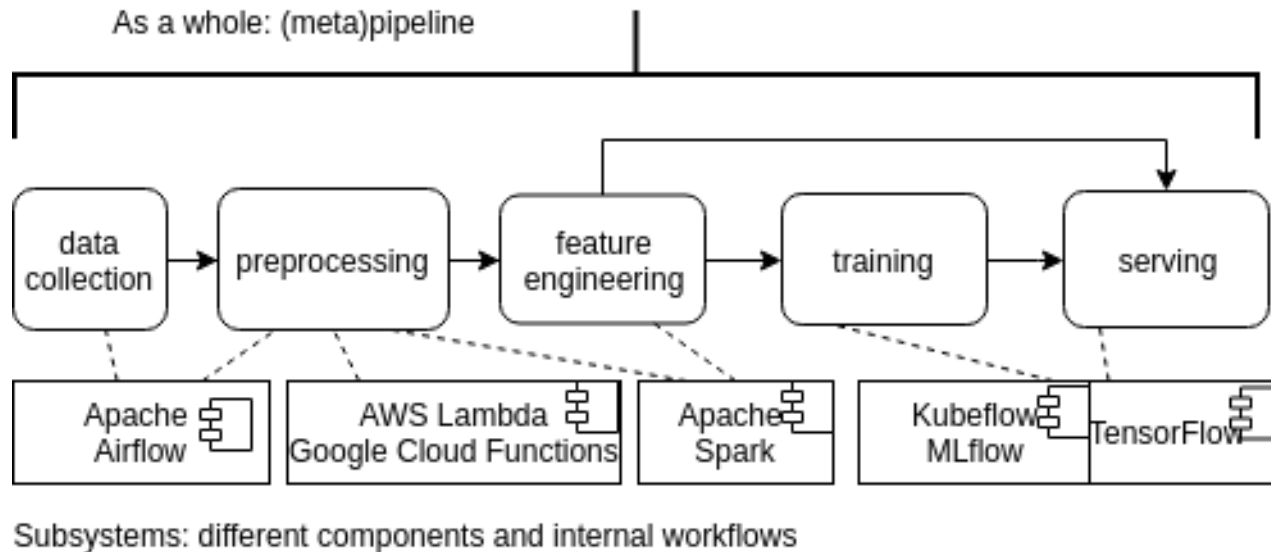


# ML systems

- **Components in machine learning**
  - machine learning algorithms is a kind of “data processing”
  - there are many other components for data-preparation, data management, experiment management
- **Machine learning pipelines**
  - complex structured components, (meta)workflows
- **Data**
  - training/validation/test data, and data to be inferenced
  - models and parameters, ML experiment settings and data
  - from the big data platforms viewpoint: they are all data!

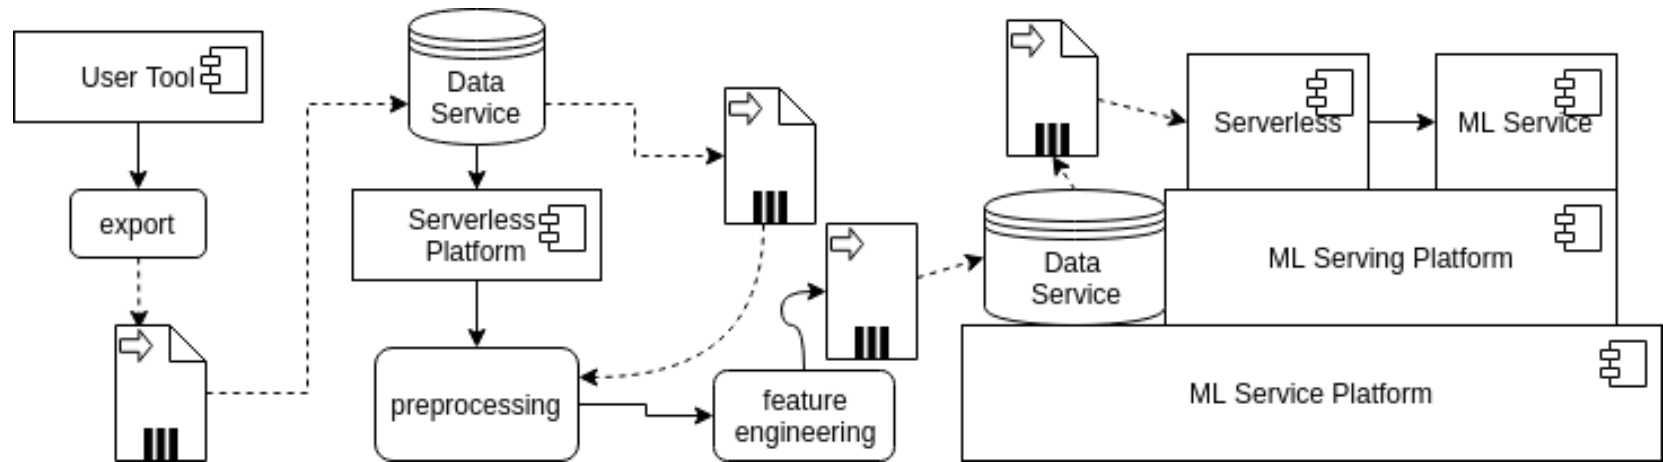
# ML workflows

- **Two possible levels:**
  - meta-workflow or pipeline
  - inside each phase: pipeline/workflow or other types of programs



# An example

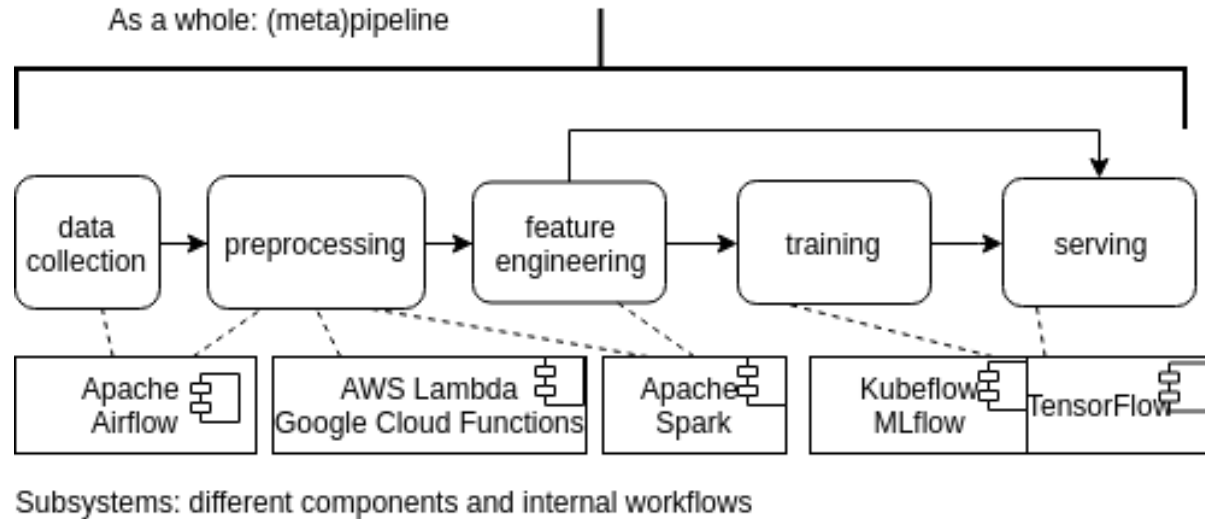
## Classifying objects in Building Information Model (BIM) in Architecture, Construction and Engineering



# System view: common characteristics of big data and ML systems?

- **(Static) system structures and functions**
  - include components, algorithms, input/output data
  - viewed as a whole, sub-systems, and individual parts
- **Computing and data infrastructures/platforms**
  - virtual machines/containers, brokers, storage, orchestration
- **Runtime quality/capability**
  - fault-tolerance, high-performance, high availability, secure, etc.

# Examples of common components in big data and ML systems



**Big data storage/ingestion**

**Big data processing**

**Resource management, workflow execution, data management tools, etc.**

# Computing and data infrastructures


# Cloud/HPC

- **Clusters of VMs/containers**
  - e.g., in Aalto we use CSC (<https://www.csc.fi/>)
- **High performance systems**
- **Known accelerators**
  - GPU and FPGA
- **New AI Accelerators/Processing Units**
  - TPU (Tensor Processing Unit)
  - Neural Network Processor (NNP)
  - Vision Processor Unit (VPU)
  - IPU( Intelligent Processing Unit)

# Edge systems

## New types of edge and edge-cloud

**Coral with Edge TPU**  
**System-on-Module, Google**  
**Edge TPU ML accelerator**  
**coprocessor**




**Jetson NVIDIA (GPU+CPU)**





# Harnessing and orchestrating end-to-end resources

End-to-end  
Edge-Cloud  
Resources



# New quantum computing for ML?

Home / Services / Azure Quantum

## Azure Quantum PREVIEW

Experience quantum impact today on Azure

[Become an early adopter](#) [Get started with the Quantum Development Kit >](#)

## Amazon Braket

Explore and experiment with quantum computing

[Sign up for the preview](#)

## Quantum Computing Playground

[quantumplayground.net](https://quantumplayground.net)



## IBM Quantum Experience

is quantum on the cloud

Accelerate your research and applications with the next generation of the leading quantum cloud services and software platform.

[Try it out now](#) [🔗](#)



Atos Quantum Learning Machine. Photo: Atos.

**Kvasi — CSC acquires quantum computing simulator**

**<https://www.csc.fi/en/-/kvasi-csc-acquires-quantum-computing-simulator>**

# Examples of common infrastructural/platform components

- **Data collection, ingestion, verification**
  - also data versioning management
- **Algorithms and serving components**
  - serving platforms and infrastructures
- **Configuration and workflow execution management**
- **Observability, monitoring and analysis**
- **Resource management and orchestration**

# Runtime abilities/capabilities

**Can you name some runtime abilities/capabilities that are important for your big data/ML systems?**

# Examples

- **Performance**
- **Accuracy**
- **Cost**
- **Scalability**
- **Failures handle/incidents management**
- **Site Reliability Engineering (SRE) concepts:**
  - Service level agreement (SLA), service level objective (SLO) and service level indicator (SLI)
    - *<https://landing.google.com/sre/sre-book/toc/index.html>*

# Robustness, Reliability, Resilience and Elasticity (R3E)

# Our objectives for end-to-end Big Data/ML systems engineering

- **Deal with end-to-end aspects that real world requires**
  - not just ML models
- **Reduce software and data engineering time**
- **Scale our systems**
  - big data, large-scale infrastructures and high number of customers
- **Optimize the system under various constraints**
- **Offer a production-level “reliable service” for customers**



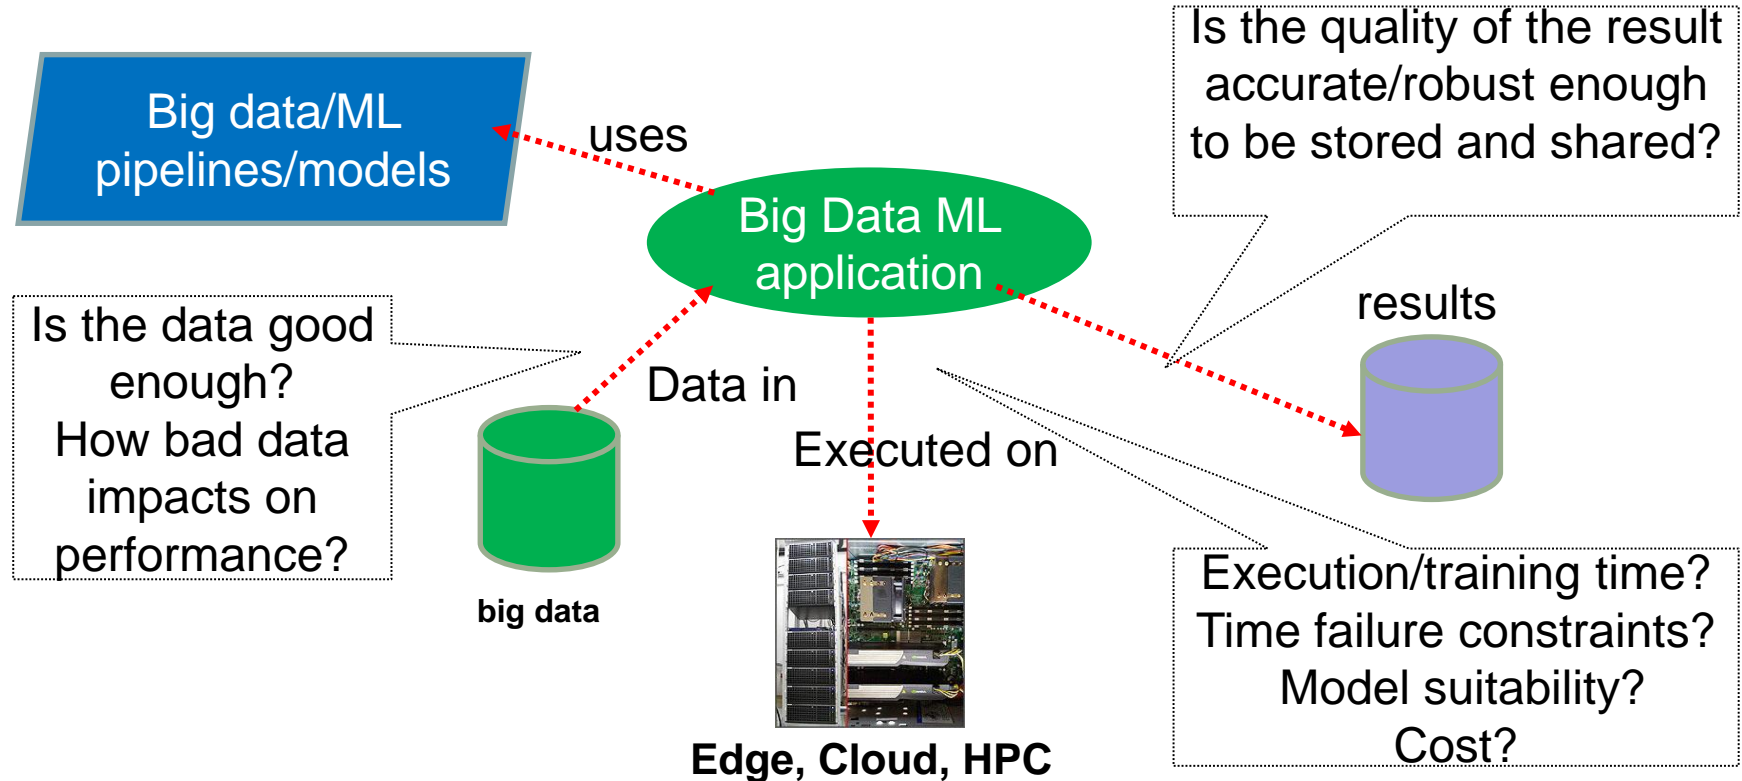
# The complexity of end-to-end view

- **Engineering, optimizing and operating big data/ML systems**
  - which are key abilities that we should define, design, monitor, and measure?
  - how do we manage software artefacts, data, configuration, ...?
  - how to enable flexibility and execution management?
  - how to prepare for “future”/”emerging” infrastructures?
  - which are tools and frameworks that help reducing engineering complexity?

# Key areas in our concerns

- **Software development**
  - testing, experimenting, benchmark, optimization, cost management
- **Resource management**
  - execution atop multiple computing frameworks suitable for ML, such as Clouds, Supercomputing, edge, ...
- **(Runtime) Ability/Quality Assurance**
  - specification, monitoring and assurance of performance, availability, costs, reliability, etc.

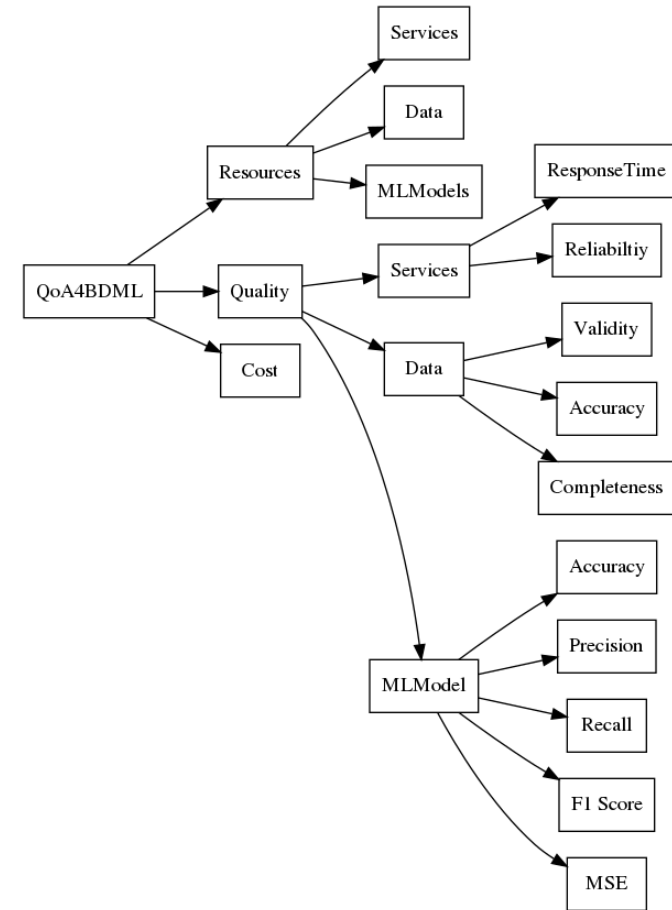
# Quality of Analytics (QoA)



**QoA = {quality of result, performance, cost}**

# Key attributes/indicators

Just example, can be more!



# Our focus – R3E

- **Robustness**
  - ability to cope with errors
- **Reliability**
  - ability to function according to the indented specification (in a proper way)
- **Resilience**
  - “ability to provide the required capability in the face of adversity”([https://www.sebokwiki.org/wiki/System\\_Resilience](https://www.sebokwiki.org/wiki/System_Resilience))
- **Elasticity**
  - ability to stretch and return to normal forms (under external forces)

# Robustness

- **In Machine Learning**
  - overfitting/underfitting
  - transfer learning
  - machine learning in an open-world
    - *how to deal with OOD (out-of-distribution) situations?*
  - when we can decide to stop training if performance/robustness does not improve?
- **In Big Data**
  - how to deal with erroneous and bad data?

# Reliability

- **System reliability versus “reliable service” (from customer/business/production view)**
- **System reliability**
  - reliable infrastructures, components, networks, ...
- **“Reliable service” → reliable data analysis/inference**
  - without failure, with specified performance
- **Some hard problems**
  - have good and enough data, clean data
  - robust pipelines without degraded performance and accuracy

# Resilience

- **Common issues in resilience**
  - distributed software and systems bugs
  - system attacks
- **Some specific issues in big data/ML systems**
  - bias in data
  - well-known problems in adversary attacks in ML phases



# Elasticity

- **Add and remove resources**
  - CPUs, memory, data, networks, ...
- **Dynamic changes of algorithms**
- **Shift computation between edge and cloud infrastructures dynamically**
  - cloud data centers, edge systems and edge-cloud systems
- **Add/remove data to improve performance**
- **Hyperparameter tuning tradeoffs**

# Short summary

Attributes	Cases from big data view	Cases from machine learning view
Robustness	deal with erroneous and bad data [45], data processing job robustness	dealing with imbalanced data, learning in an open-world (out of distribution) situations [23, 34, 35]
Reliability	reliable data sources, support of quality of data [28, 46], reliable data services [26], reliable data processing workflows/tasks [47]	reliable learning and reliable inference in terms of accuracy and reproducibility of ML models [22, 34]; uncertainties/-confidence in inferences; reliable ML service serving
Resilience	software bugs, infrastructural resource failures, fault-tolerance and replication for data services and processing [44]	bias in data, adversary attacks in ML [25], resilience learning [14], computational Byzantine failures [8]
Elasticity	utilizing different data resources, increasing and decreasing data usage w.r.t. volume, velocity, quality; elasticity of underlying resources for data processing [42]	elasticity of resources for computing [19, 21, 24], elasticity of model parameters; performance loss versus model accuracy; elastic model services for performance

**Table 1: R3E with big data and ML concerns**

Source: [https://www.researchgate.net/publication/341762862\\_R3E\\_-\\_An\\_Approach\\_to\\_Robustness\\_Reliability\\_Resilience\\_and\\_Elasticity\\_Engineering\\_for\\_End-to-End\\_Machine\\_Learning\\_Systems](https://www.researchgate.net/publication/341762862_R3E_-_An_Approach_to_Robustness_Reliability_Resilience_and_Elasticity_Engineering_for_End-to-End_Machine_Learning_Systems)

# Do we need to treat

***Robustness, Reliability, Resilience  
and Elasticity***

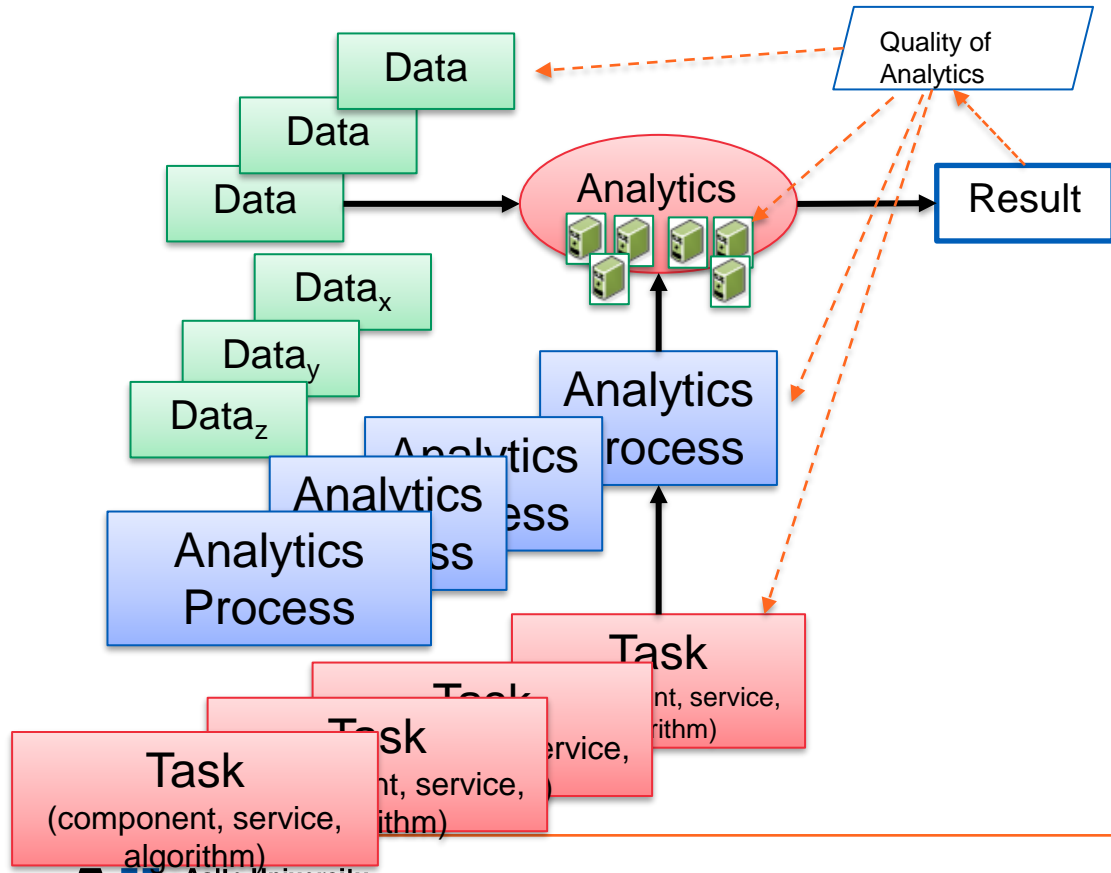
**equally in all your design?  
from which views?**

# An Approach with Elasticity Principles for R3E

# Elasticity

- **Demand elasticity**
  - elastic demands from consumers
- **Output elasticity**
  - multiple outputs with different price, quantity and quality
- **Input elasticity**
  - elastic data inputs, e.g., deal with increasing data sources
- **Elastic pricing and quality models associated resources**
  - CPU/GPU, memory/disk, networks, etc.

# Elasticity in (big) data analytics

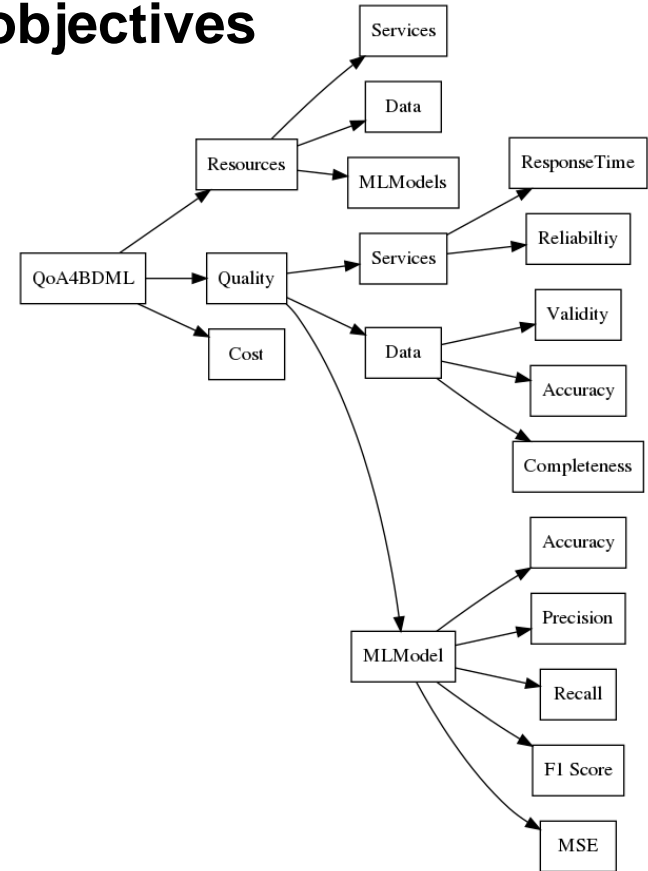


- **More data → more compute resources (e.g. more VMs)**
- **More types of data → more, different tasks → more analytics processes**
- **Change quality of analytics**
  - Change quality of data
  - Change response time
  - Change cost
  - Change types of result (form of the data output, e.g. tree, table, story)


# Establish quality of analytics for Big Data/ML

## Possible indicators/objectives

- Have clear indicators/objectives so we can establish SLA for Quality of Analytics
- You can build your own dimensions



# Elasticity engineering

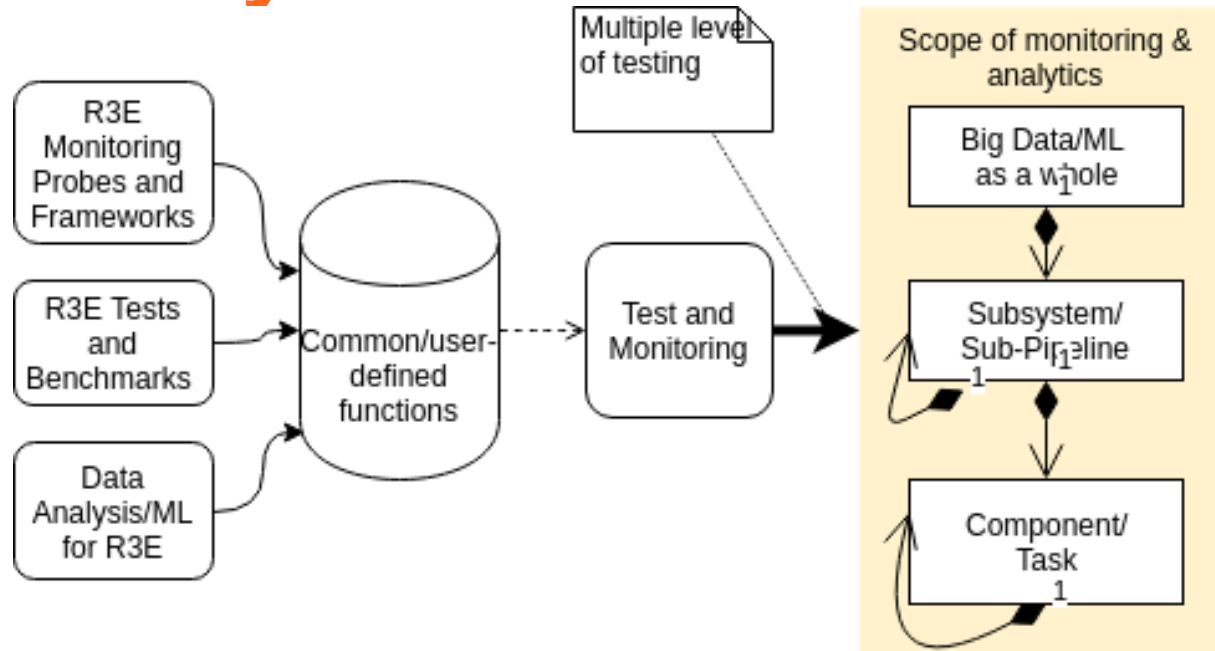




# Elasticity engineering for ML


- **Conceptualizing and modeling elastic objects**
  - ML models, computing resources, data and QoA metrics
- **Defining and capturing elasticity primitive operations**
  - change resources, QoA metrics, model parameters, input data
- **Programming features for elastic objects**
  - with ML flows, coordinating QoA adjustment, dynamic serving models
- **Runtime deploying, control, and monitoring techniques for elastic objects**

# Multi-level cross platforms monitoring and analysis



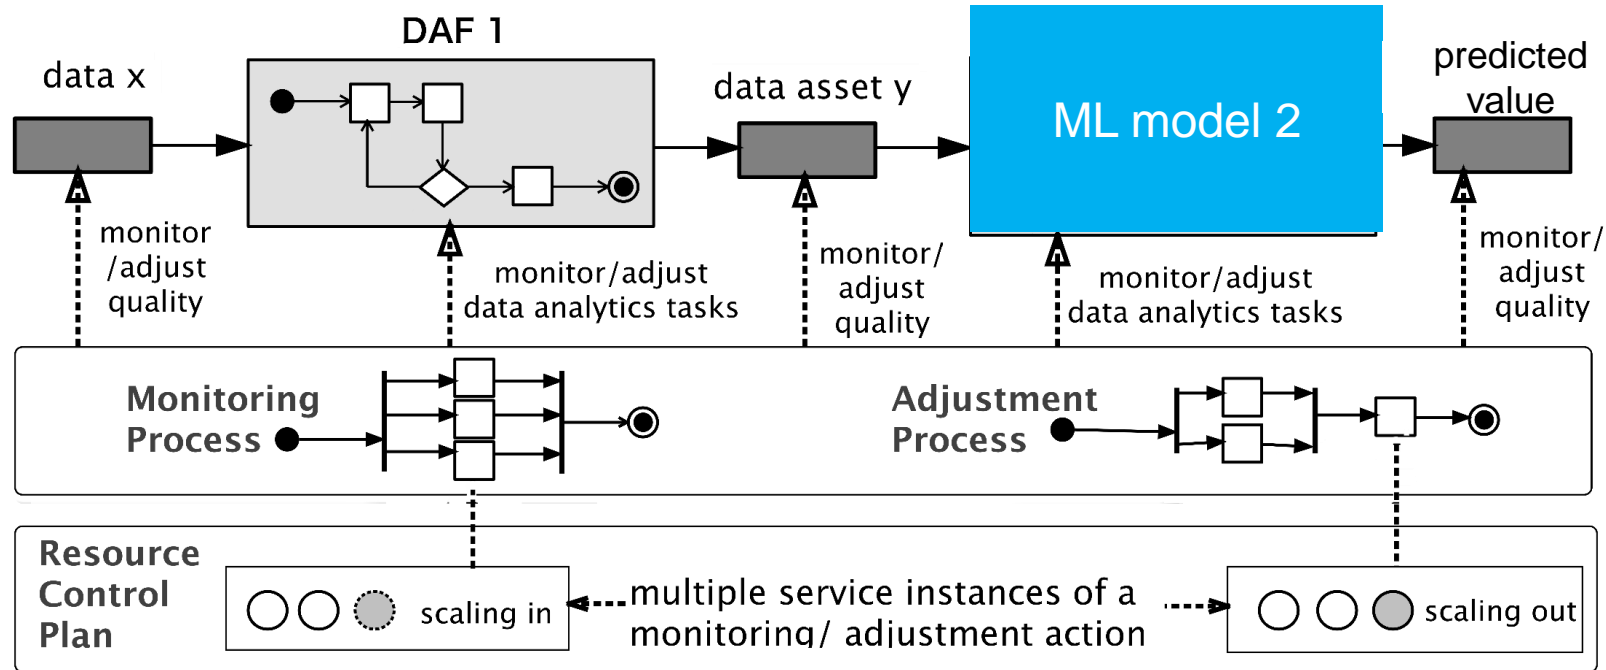
**We will have a hands-on on observability and monitoring**

# Detecting elasticity



Alessio Gambi, Daniel Moldovan, Georgiana Copil, Hong Linh Truong, Schahram Dustdar: On estimating actuation delays in elastic computing systems. SEAMS 2013: 33-42

# Using control process to ensure QoA



**Will be covered in the hands-on on elastic ML serving**

# Some examples/results

## With results from:


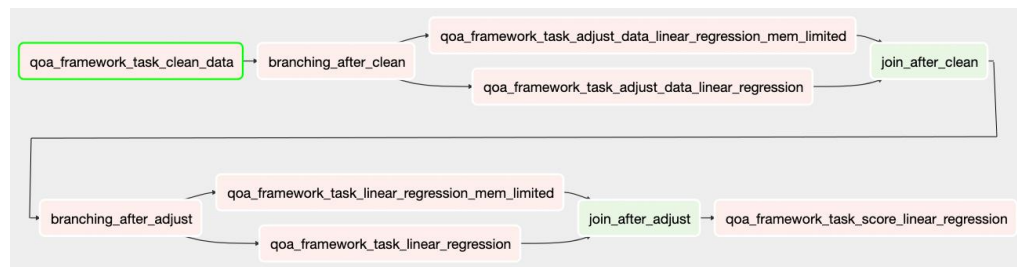
- Kreics Kristis, „*Quality of analytics management of data pipelines for retail forecasting*“, Aalto Master thesis, 2019, <https://aaltodoc.aalto.fi/handle/123456789/39908>
- Minjung Ryu, „*Machine Learning-based Classification System for Building Information Models*“, Aalto Master thesis, 2020
- Minjung Ryu, Linh Truong, Matti Kannala „*Understanding Quality of Analytics Tradeoffs in an End-to-End Machine Learning-based Classification System for Building Information Modeling*“, 2020, Working paper.
- Matt Baughman, Nifesh Chakubaji, Hong-Linh Truong, Kristis Kreics, Kyle Chard, Ian Foster, *Measuring, Quantifying, and Predicting the Cost-Accuracy Tradeoff*, IEEE International Workshop on Benchmarking, Performance Tuning and Optimization for Big Data Applications, IEEE BigData 2019, <https://research.aalto.fi/files/38801332/paper.pdf>

# Industrial retail forecast (with Sellforte)

## Forecast where to put marketing information, example of data

date	id	name	volume	price	cost	promo	category_net	margin	category1	category2	location	sales
07/01/2018	100	Chicken	38144.0	3.79	2.7	0	451692.0	0.25	Meat	Food	Helsinki	144565.76
14/01/2018	100	Chicken	36420.0	3.79	2.66	0	414342.0	0.25	Meat	Food	Helsinki	138031.8
21/01/2018	100	Chicken	35322.0	3.79	2.66	0	381854.0	0.25	Meat	Food	Helsinki	133870.38

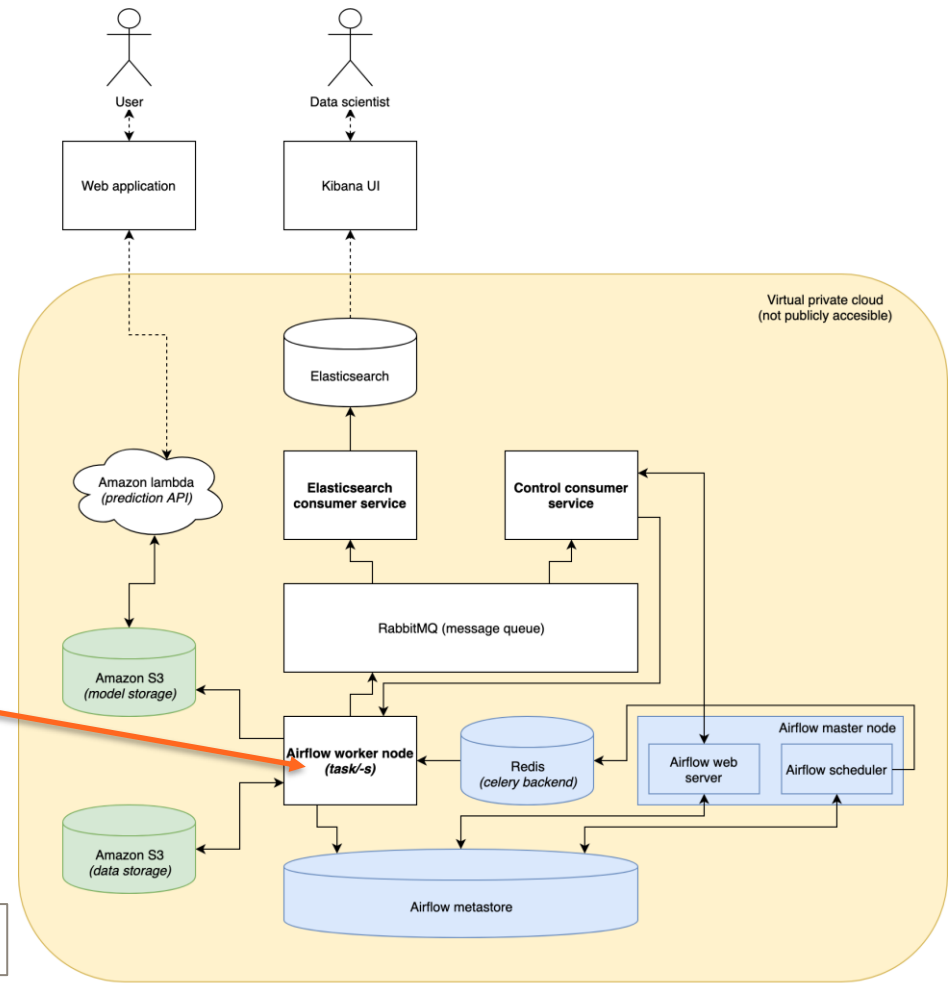
- **Metrics:**
  - data size, R square value, time, and cost
- **Pipelines**
  - tune pipelines with QoA primitive actions



Source: Kreics Krists, „Quality of analytics management of data pipelines for retail forecasting“, Aalto CS Master thesis, 2019

# Industrial retail forecast (with Selfforte)

Monitoring various indicators, including user-defined quality of data



Source: Kreics Kristis, „Quality of analytics management of data pipelines for retail forecasting“, Aalto CS Master thesis, 2019

# Initial results

## Custom cost function

```
def get_fargate_metrics_object(cpu, ram, elapsed_time, previous_result):  
    # Fargate service cost per second  
    FARGATE_CPU_COST = 0.04048 / 60 / 60  
    FARGATE_RAM_COST = 0.004445 / 60 / 60  
    if previous_result and 'cost_usd' in previous_result:  
        cpu_cost = previous_result['cost_cpu'] + FARGATE_CPU_COST  
        ram_cost = previous_result['cost_ram'] + (ram['used']/1024/1024/1024) * FARGATE_RAM_COST  
    else:  
        cpu_cost = FARGATE_CPU_COST  
        ram_cost = (ram['used']/1024/1024/1024) * FARGATE_RAM_COST  
    return { 'cost_cpu': cpu_cost, 'cost_ram': ram_cost, 'cost_usd': ram_cost + cpu_cost }
```

## Custom instrumentation for model quality

```
# model_score returns a dict -> { 'r2_squared': r2_squared_score }  
model_score = score_model(store, model, data_path, preset)  
pm.log_analytics_metric(model_score)
```

Source: Kreics Kristi, „Quality of analytics management of data pipelines for retail forecasting“, Aalto CS Master thesis, 2019

## Examples of **actions** in Elasticity Primitive Operations

```
def default_get_control_action(body_dict):  
    index = body_dict.pop('metric_type', None)  
    print(body_dict, flush=True)  
    try:  
        if index == 'metrics':  
            if body_dict['cost_usd'] > 1 or body_dict['time_elapsed'] > 500:  
                return 'SOFT_STOP'  
            elif body_dict['time_elapsed'] > 1000:  
                return 'HARD_STOP'  
        elif index == 'data_logs':  
            if body_dict['task_name'] == 'clean_data':  
                if body_dict['in']['train.csv'] / 2 > body_dict['out']['train.csv']:  
                    return 'SOFT_STOP'  
        elif index == 'analytics':  
            if body_dict['payload']['r2_squared'] < 0.2:  
                return 'SOFT_STOP'  
    else:  
        print('No valid index found!')  
        return -1  
    except KeyError:  
        pass
```



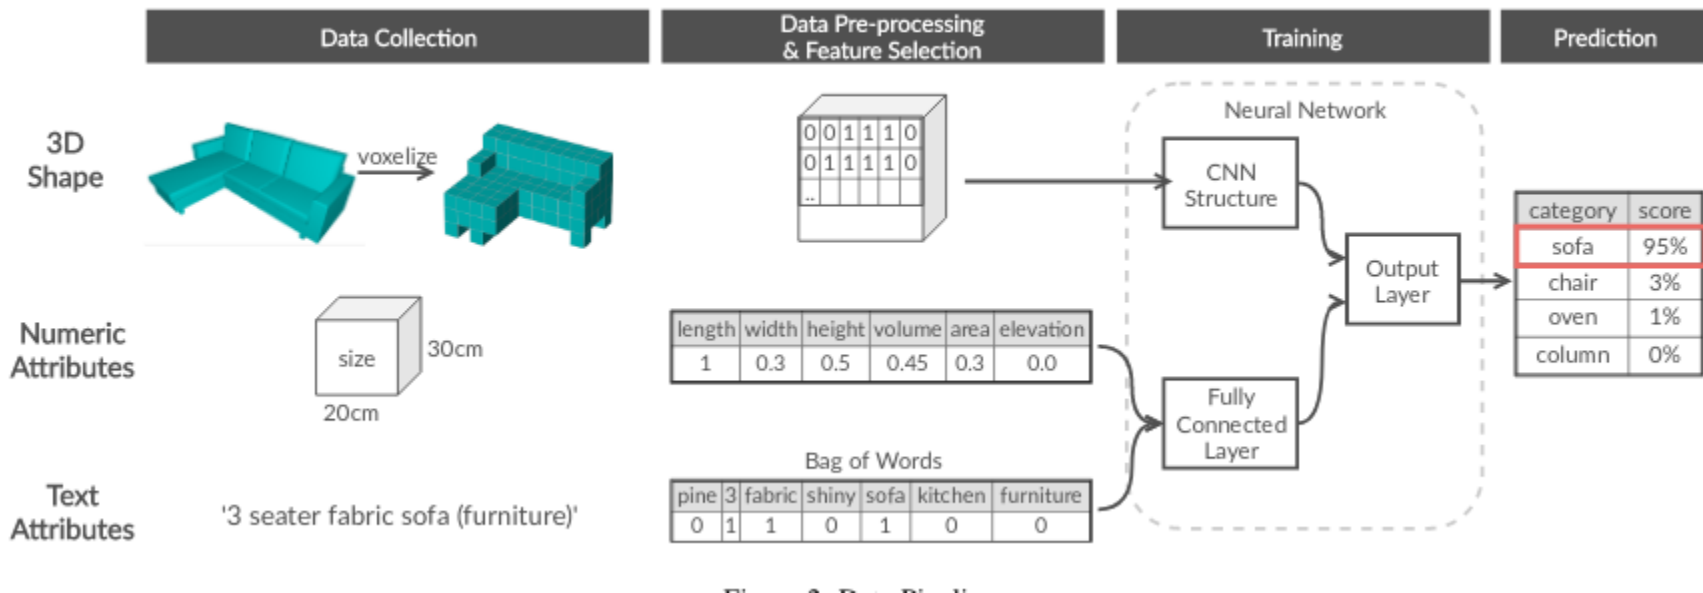
# Initial results

- Running with Airflows in Amazon EC2
- Apply different actions to change “store” (domain objects) and computing resources
- Real improvement (from the domain expert) with 1 million rows case

13.3% lower accuracy and 44% shorter time, R squared value was 9.5% lower → could good enough results for 50% of total store locations

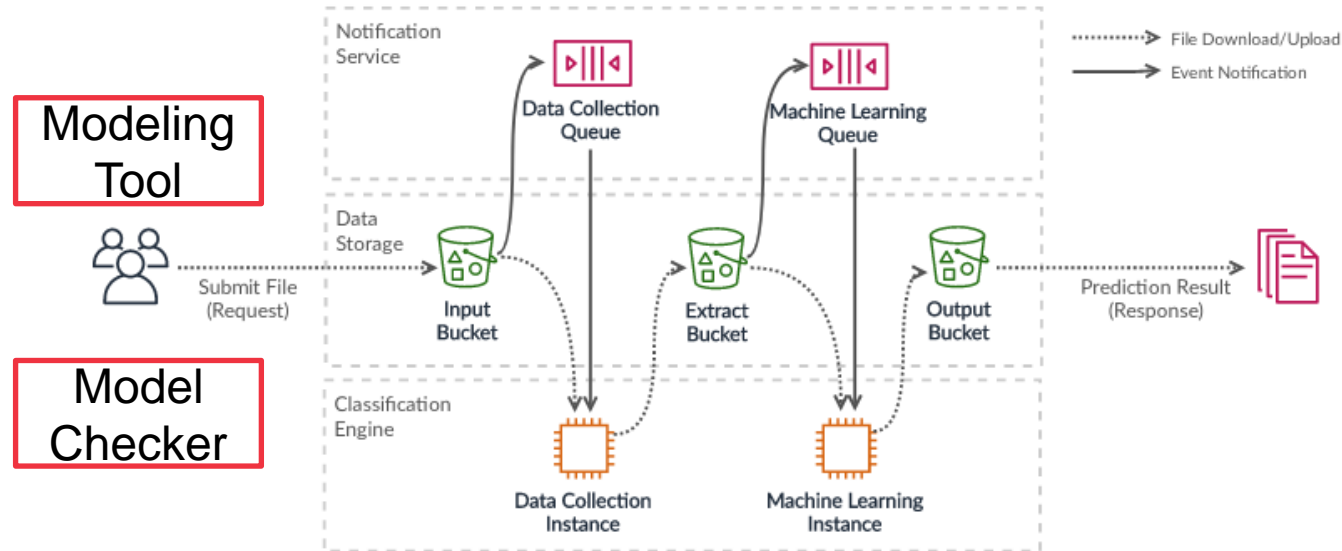
The application-aware data reduction strategy and cost-accuracy tradeoffs may be more intelligently made based on knowledge of the application domain.

# ML classification for BIM (with Solibri data)



Source: Minjung Ryu, „Machine Learning-based Classification System for Building Information Models“, Aalto CS Master thesis, 2020

# ML classification for BIM (with Solibri data)



Source: Minjung Ryu, „Machine Learning-based Classification System for Building Information Models “, Aalto CS Master thesis, 2020

# Initial results

- Data set: 591 classification cases from 146 models
- Machines: AWS/Local with/out GPUs
- Different cases and settings






Figure 5: Impact of object counts on DC time and on ML time


**Reveal various relationships between types of data, extracting data resolution, machines and the accuracy of classifications**



(a) Dim vs DC time



(b) Dim vs ML time



(c) Dim vs Accuracy

# Study log for this week

## Think about

- What does it mean R3E for *YOUR big data and machine learning systems*?

## Then

- in your experience/work, which ones of R3E concern you most? Why? What would you do? What do you look for?
- ~1 page – submit into the Mycourses for comments/feedback (keep it in your git)

# Thanks!

**Hong-Linh Truong**  
**Department of Computer Science**

**rdsea.github.io**