



Aalto University
School of Science

Coordination Models and Techniques for Big Data and Machine Learning Systems

Hong-Linh Truong
Department of Computer Science
linh.truong@aalto.fi, <https://rdsea.github.io>

Learning objectives

- **Analyze the role of coordination techniques, their complexity and diversity in big data/ML systems**
- **Understand and apply orchestration models, common tools and design patterns**
- **Understand and apply choreography models, common tools and design patterns**
- **Understand, define and develop ML Model Serving**

Coordination complexity and diversity

Examples of common tasks

Discussion:

- ML phases & tasks
- Software stack
- Execution environments
 - Computing resources
- R3E

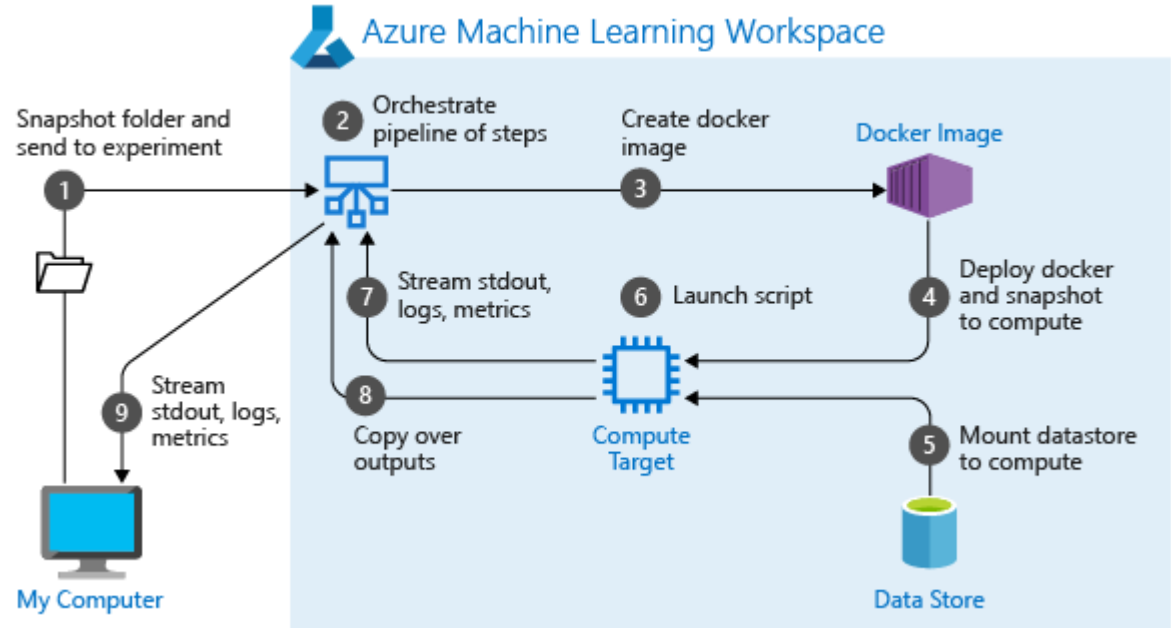
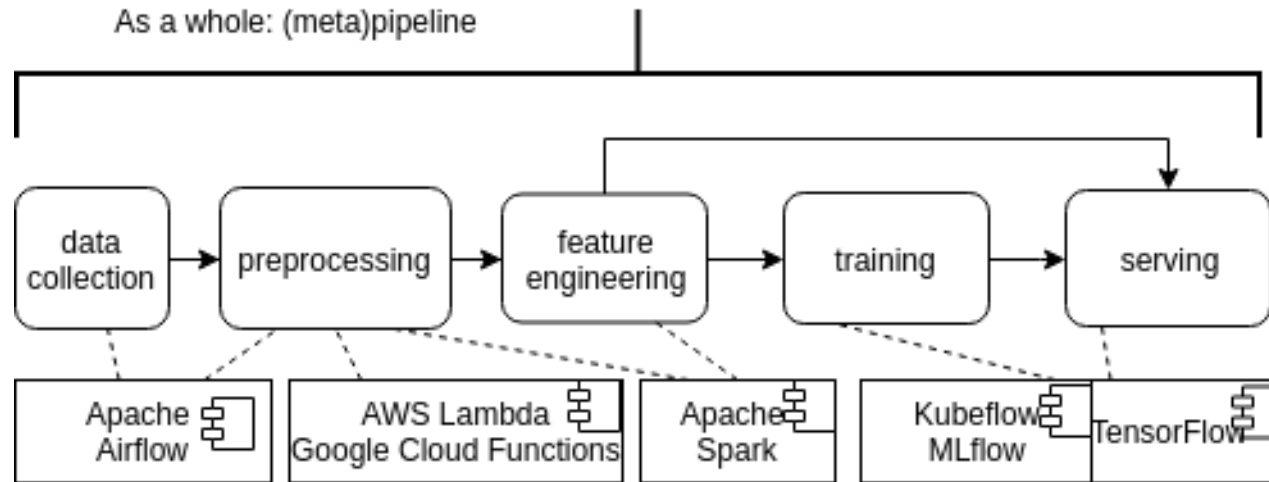


Figure source: <https://docs.microsoft.com/en-us/azure/machine-learning/concept-ml-pipelines>

Recall: big data/ML systems

- **Multiple levels:**
 - Meta-workflow or -pipeline
 - Inside each phase: pipeline/workflow or other types of programs



**Automation vs
manual tasks:
share your
current way?**

Main issues related to coordination

- **How to coordinate phases and tasks in big data/ML systems**
 - automation is an important requirement, why?
- **How to prepare artefacts and resources for big data/ML systems**
- **How to manage tests and experiments**
 - trial computing configurations, inputs/results collection
- **How to control for assuring R3E for the pipeline execution**
 - end-to-end R3E requires coordination
 - issues in internal and external services

W3H: what, when, where and how for coordination

Where: within a phase, across phases, within a component, a subsystem, etc.

What: preparing data and machines, performing inferences, carrying out observability

When: triggered by data flows or control flows or messages/events?

How: which tools, models?

Coordination



Diversity and Complexity

- **Diversity**

- so many tools/frameworks in a single big data/ML system
→ *a single coordination model/tool might not be enough*
- there exist many coordination systems (included your specific implementation)
→ *which ones should we select?*

- **Complexity, due to**

- integration models with big data/ML components and infrastructures
- very large-scale scale
- runtime management: performance, failures, states

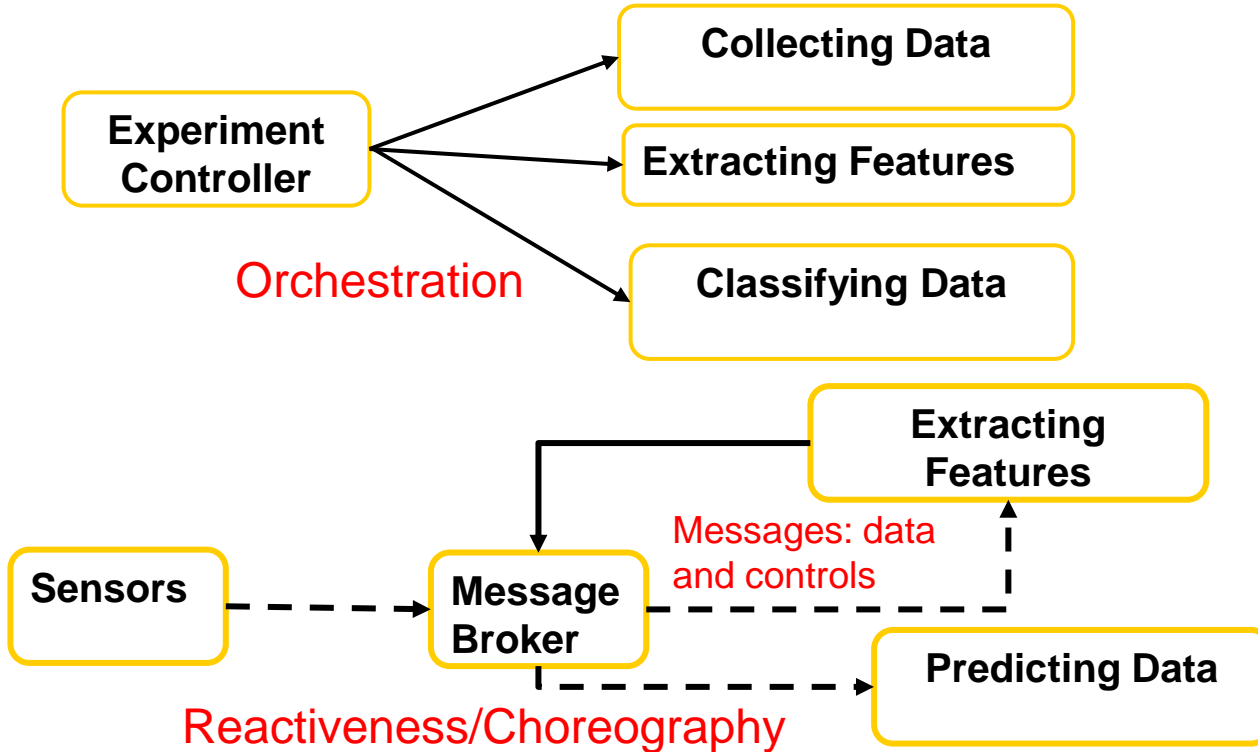
Key notions

- **Workflow and Task/Activity/Step abstraction**
 - a task can encapsulate a “complex workflow”
- **Big Data/ML software frameworks**
 - for implementing big data/ML capabilities
- **Platform services for coordination**
 - services offering features/functionality for executing “tasks”
 - single or multiple providers?
- **Execution environments and resources**
 - single platform or cross (heterogeneous) platforms

Coordination styles

- **Coordination models for Big Data/ML systems**
 - orchestration and reactivity/choreography
- **Orchestration**
 - task graphs and dependencies are based on control or data flows
 - dedicated orchestrator: tasks triggered based on completeness of tasks or the availability of data
- **Reactivity/choreography**
 - follow reactive model: tasks are reacted/triggered based on messages

Orchestration and Reactiveness



System issues impacting coordination

- **Main situations:**
 - within the same system/infrastructure
 - *all services and computing resources belong to the same platform/infrastructure*
 - *e.g., running everything with Google Cloud or Microsoft Azure*
 - across systems/infrastructures
 - *services in different clouds or cloud data centers*
 - *e.g., Edge-cloud infrastructures*
 - with the same software stack or not?
- **How such situations would affect the coordination?**



Aalto University
School of Science

Coordination with workflow techniques

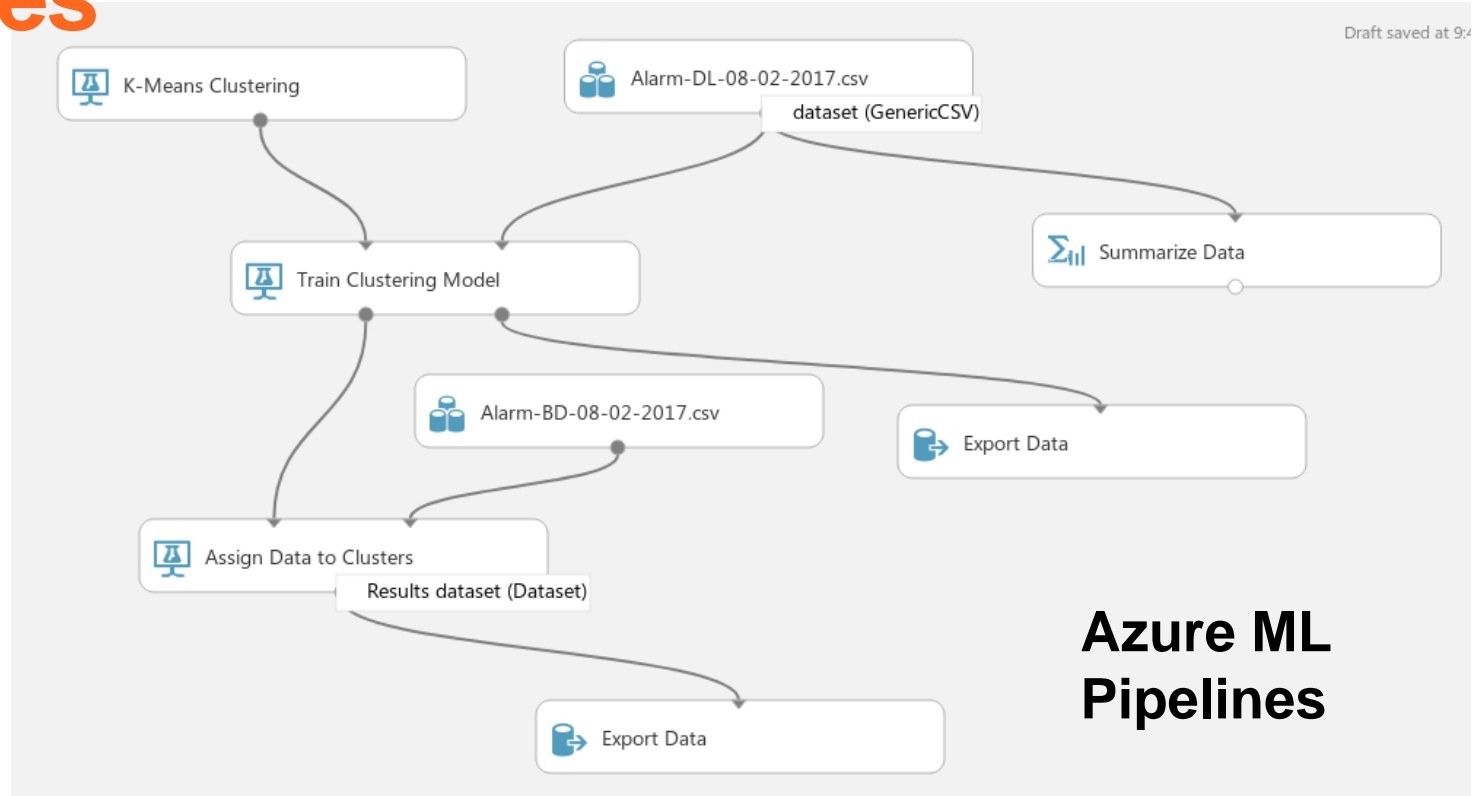
The orchestration style

Orchestration architectural style: design

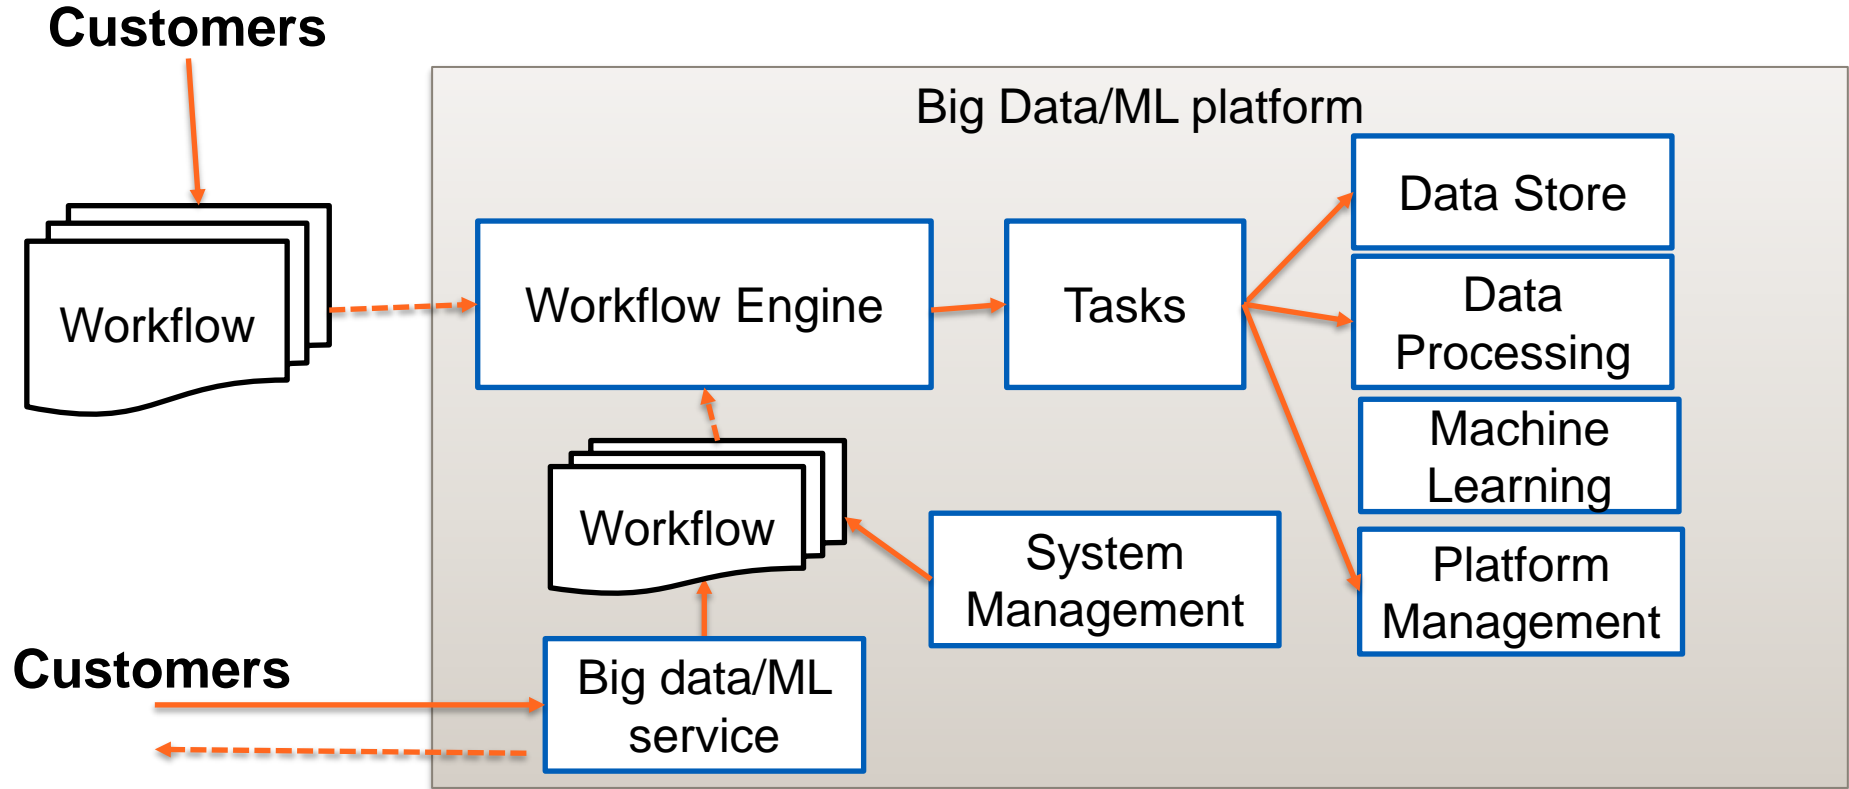
- **Workflow architectures are known**
 - Big Data/ML systems: leverage many types of services and cloud technologies
- **Required components**
 - workflow/pipeline specifications/languages (also UI)
 - data and computing resource management
 - orchestration engines (with different types of schedulers)
- **Execution environments**
 - cloud platforms (e.g., VMs, containers, Kubernetes)
 - heterogenous computing resources (PC, servers, Raspberry PI, etc.)

Example: workflow used in ML pipelines

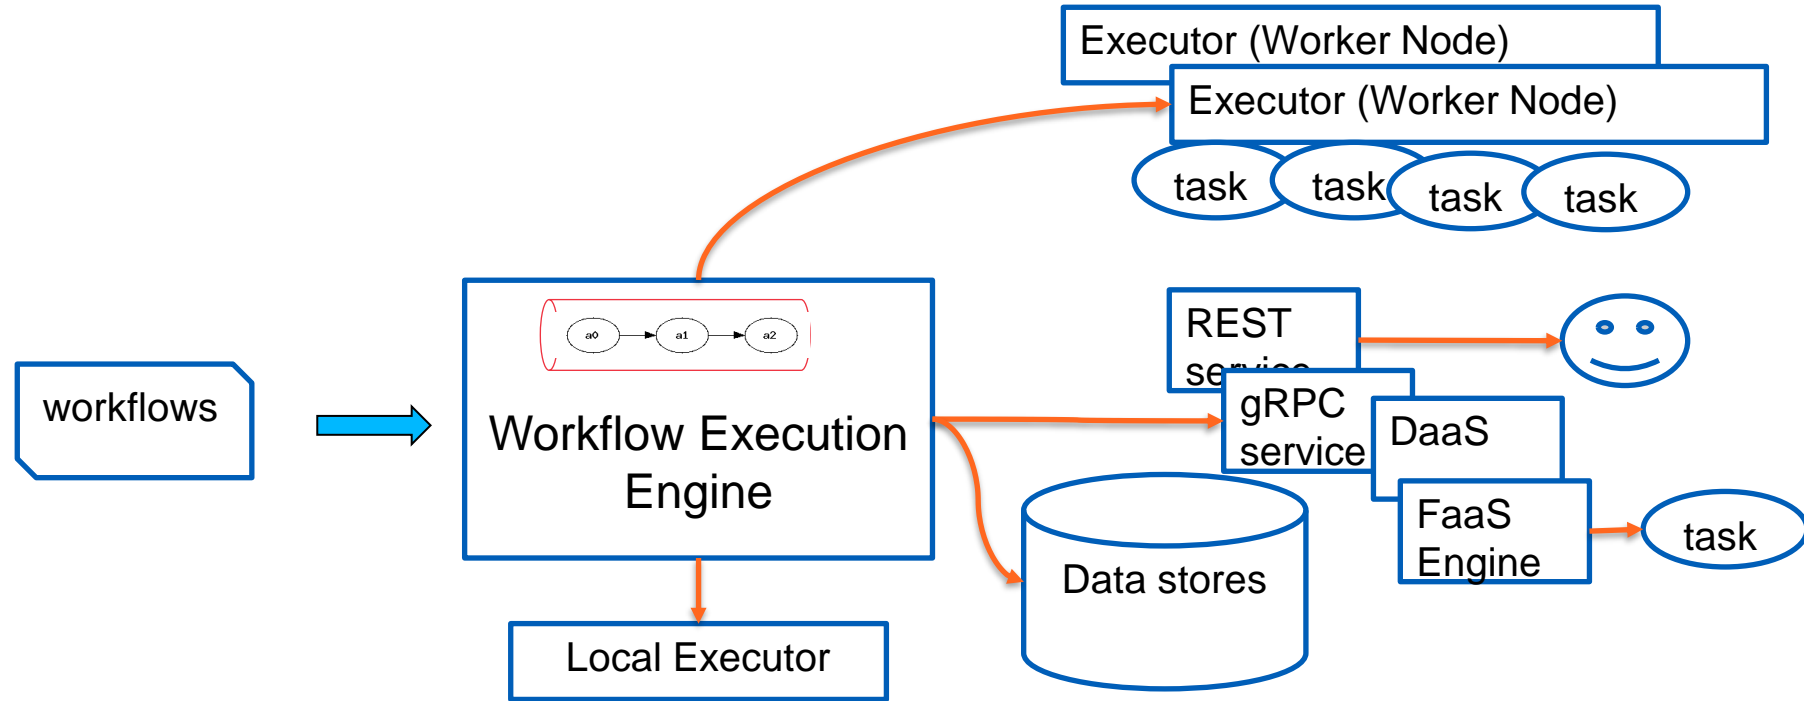
So what is behind the scene?



Workflows in big data/ML systems



Common workflow execution models



Executors: containers, VMs, common OS processes

Key components

- **Tasks/Activities**

- describe a single work (it does not mean small)
- tasks can be carried out by humans, executables, scripts, batch applications, stream applications and services.

- **Workflow Languages**

- how to structure/describe tasks, dataflows, and control flows

- **Workflow Engine**

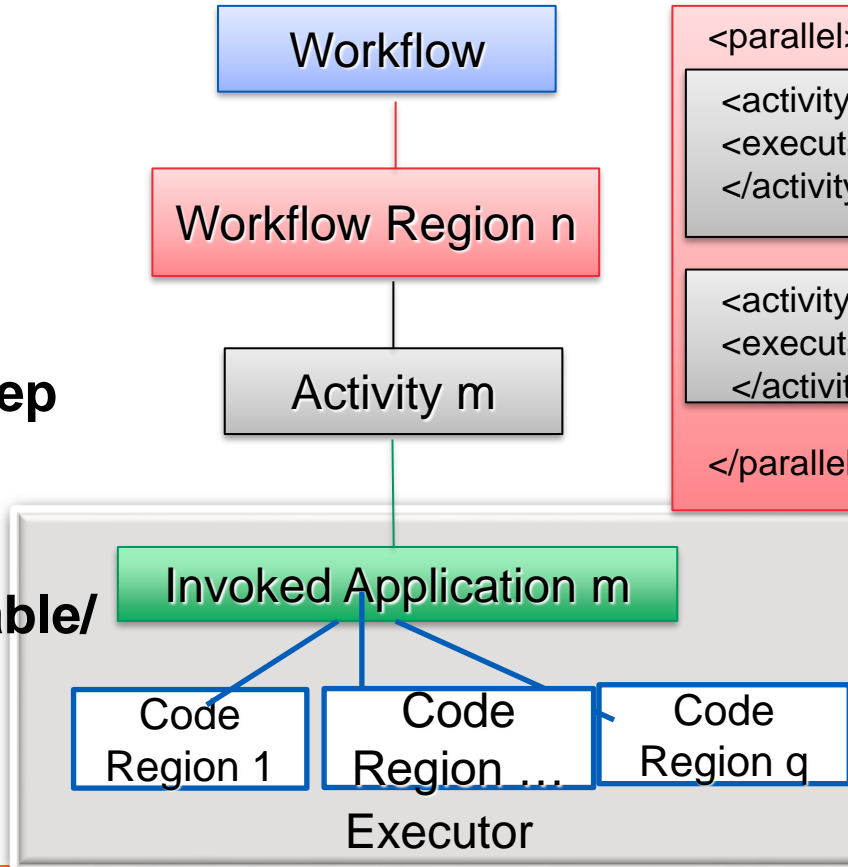
- execute the workflow by orchestrating tasks
- usually call remote services to run tasks

Structured view of workflows

abstract example

Task/Step

Service
/executable/
script



<parallel>

```
<activity name="mTask1">
<executable name="mTask1"/>
</activity>
```

```
<activity name="pTask2">
<executable name="cbpprediction"/>
</activity>
```

</parallel>

```
cbpprediction.py
def find_target_regression_model(dataset):
```

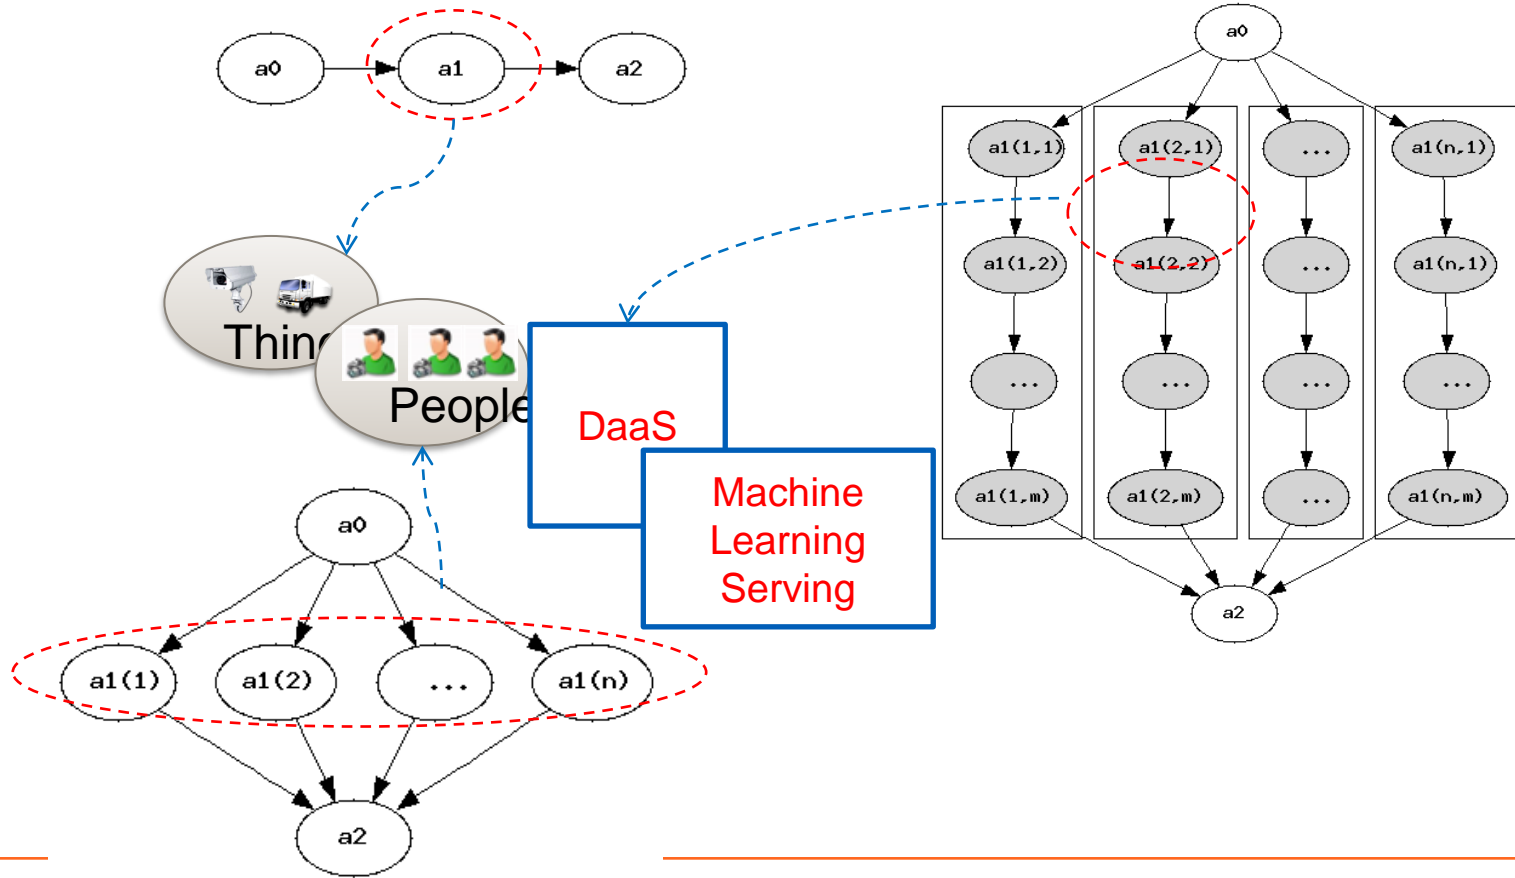
```
    selecteddata=selecteddata[TARGE
    T_FEATURES].dropna()
```

```
    target_lin_reg =LinearRegression()
    target_lin_reg.fit(X,Y)
```

Structured view

- **Invoked applications within a task/activity**
 - can be a containerized service, script, python program, a function, etc.
 - can be designed for different purposes: e.g. computation, management or prediction
- **Encapsulating the whole workflow**
 - the whole workflow can be encapsulated with in a service
 - thus the whole workflow can be invoked via a service call for multiple consumers

Tasks orchestration



Runtime aspects

- **Parallel and distributed execution**
 - tasks are deployed and running in different machines
 - multiple workflows are running
- **Long running**
 - can be hours!
- **Checkpoint and recovery**
- **Monitoring and tracking**
 - which tasks are running, where are they?
- **Stateful management**
 - dependencies among tasks w.r.t control and data

Describing workflows

- **Programming languages with procedural code**
 - general- and specific-purpose programming languages, such as Java, Python, Swift
 - common ways in big data and ML platforms
- **Descriptive languages with declarative schemas**
 - BPEL, YAML, and several languages designed for specific workflow engines
 - common in business and scientific workflows
 - YAML is also popular for big data/ML workflows in native cloud environments

Workflow frameworks

- **Often running in the same infrastructure**
- **Task-driven or data-driven specification**
- **Generic workflows**
 - Use to implement different tasks, such as machine provisioning, service calls, data retrieval
 - *Examples: Airflow, Argo Workflows*
- **Specific workflows for specific purposes**
 - E.g., Kubeflow (<https://github.com/kubeflow/pipelines>)

Workflow frameworks

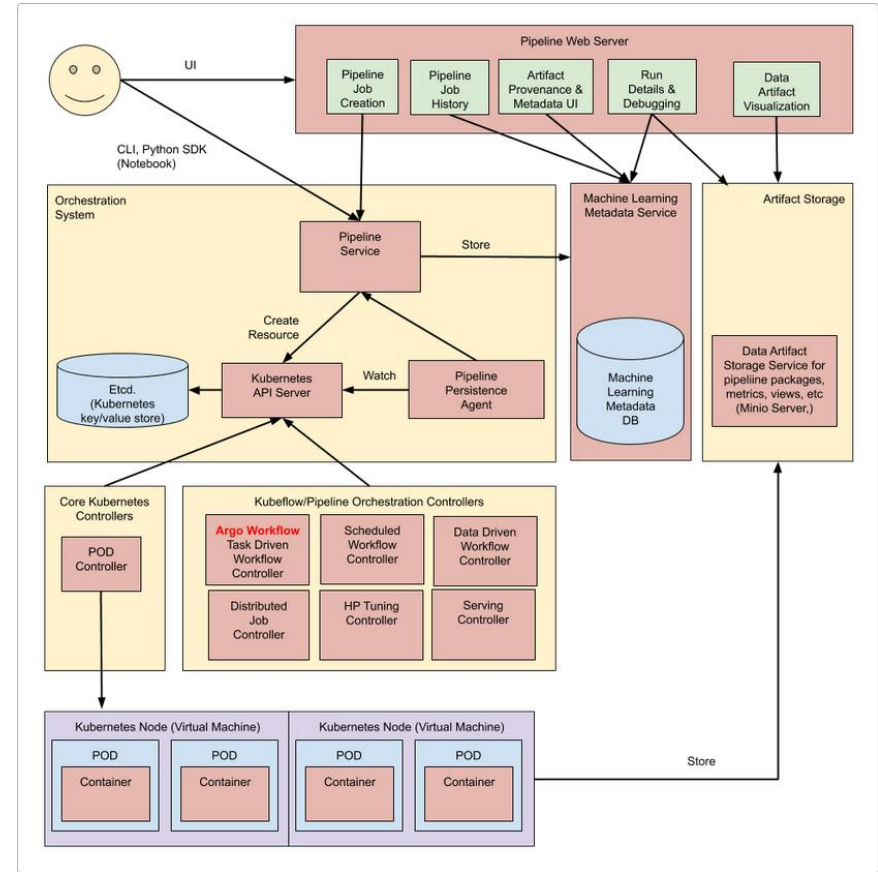
- **Apache Airflow**
 - <https://airflow.apache.org/>, also used by Google Composer
- **Argo**
 - <https://argoproj.github.io/argo/>, used in Kubeflow
- **Prefect**
 - <https://www.prefect.io/>
- **Uber Cadence**
 - <https://github.com/uber/cadence>

What purposes are for using workflows?

- **For coordinating big data/ML phases/stages in the pipeline**
- **For implementing tasks within a phase/stage**
 - implement data preprocessing task
 - implement training tasks
 - implement experiment management
 - batch model for machine learning serving
- **Which ones? No easy answer**
 - separate – as a framework/service
 - integrated within big data/ML frameworks

Examples: Kubeflow

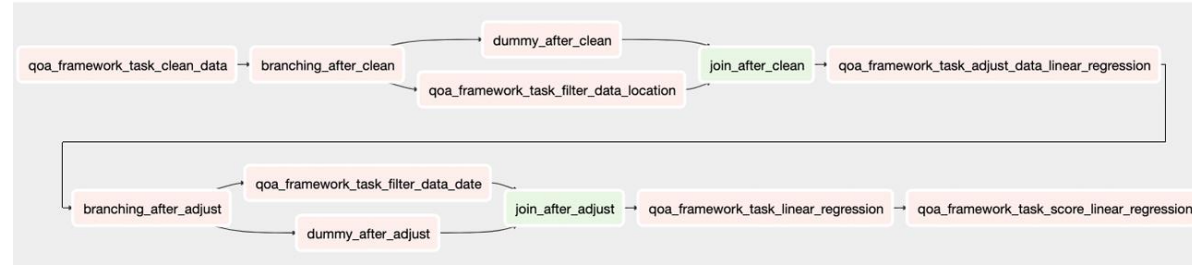
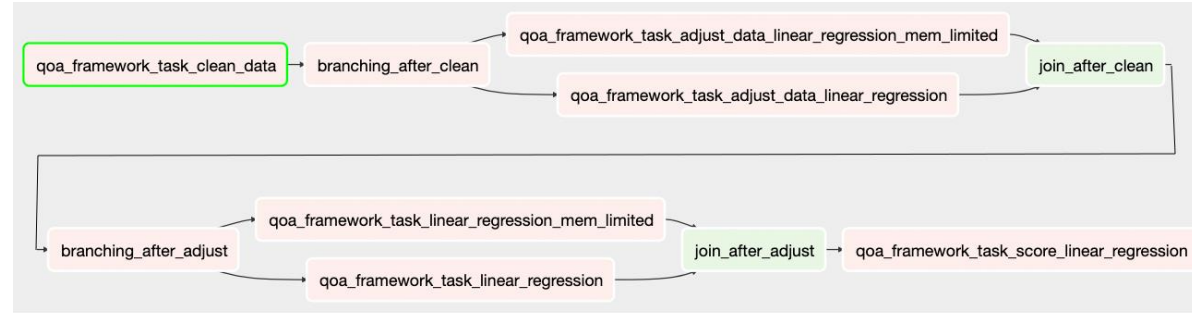
- End-to-end Orchestration
 - Orchestration is based on workflows
 - Using “Orchestration controllers”
- Discussion: dealing R3E with ML workflows?
 - Where, What, When and How



<https://www.kubeflow.org/docs/pipelines/overview/pipelines-overview/>

Examples: Coordinating tasks

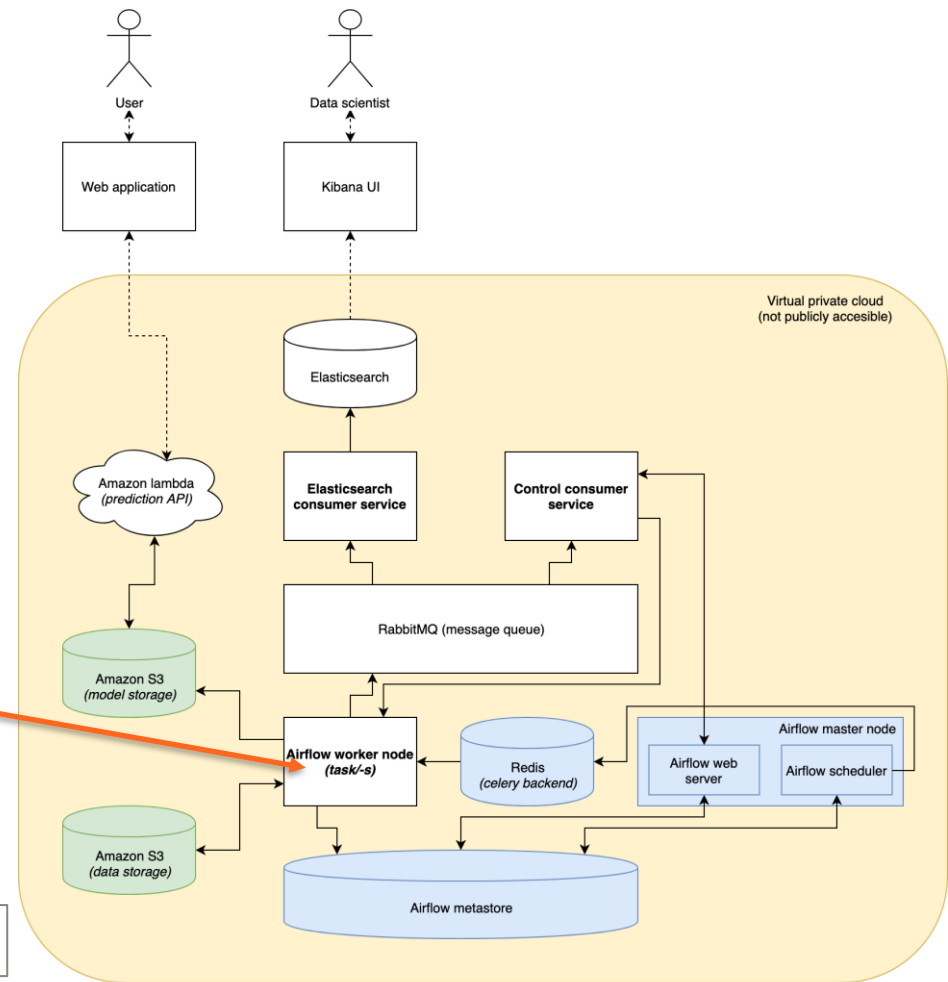
- Discussion: dealing R3E with ML workflows?
- Where, What, When and How



Source: Kreics Kristis, „Quality of analytics management of data pipelines for retail forecasting,“, Aalto CS Master thesis, 2019

Examples: Exchanging metrics for R3E coordination

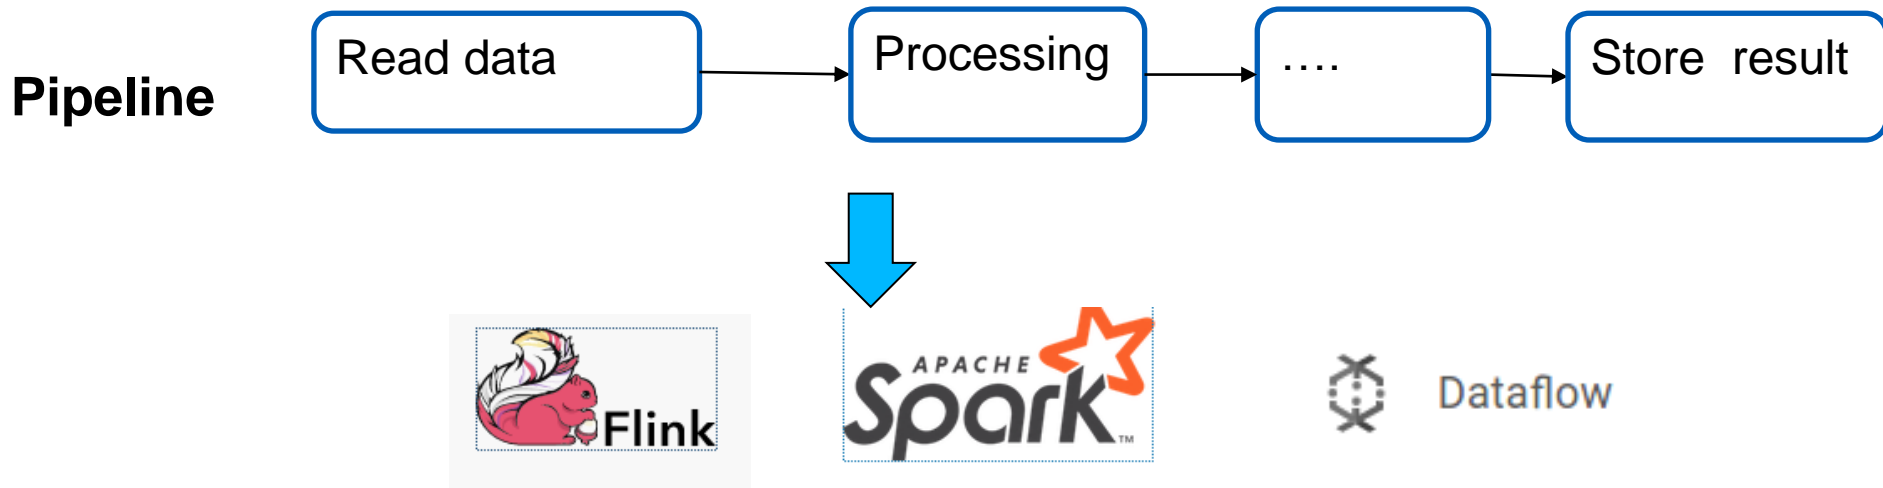
Monitoring various metrics,
including user-defined
quality of data



Source: Kreics Krists, „Quality of analytics management of data pipelines for retail forecasting“, Aalto CS Master thesis, 2019

Example: Apache Beam

- **Goal: separate from data processing pipelines from backend execution engines**
 - Focus on pipeline design



Example: Apache Beam

- <https://beam.apache.org/>
- **Suitable for data analysis processes that can be divided into different independent tasks**
 - ETL (Extract, Transform and Load) & Data Integration
 - ML pipeline implementation
- **Execution principles:**
 - Mapping tasks in the pipeline to concrete tasks that are supported by the selected back-end engine
 - Coordinating task execution like workflows.

Example: Apache Beam

Basic programming constructs

- **Pipeline:**
 - For creating a pipeline
- **PCollection**
 - Represent a distributed dataset
- **Transform**
 - $[\text{Output PCollection}] = [\text{Input PCollection}] \mid [\text{Transform}]$
 - Possible transforms: ParDo, GroupByKey, Combine, etc.
 - Partition: split the data

Example: Apache Beam

- **Data preprocessing and featuring engineering**
 - could also be for data validation
- **Preparation for training**
 - processing and partitioning data
- **Inferences**
 - implement inference functions which can be called within a pipeline, e.g. using ParDo()


Coordination with messaging

The reactiveness/choreography

Choreography: Reactive systems for Big Data/ML

Do you remember key principles of reactive systems?

Reactive systems



Source: <https://www.reactivemanifesto.org/>

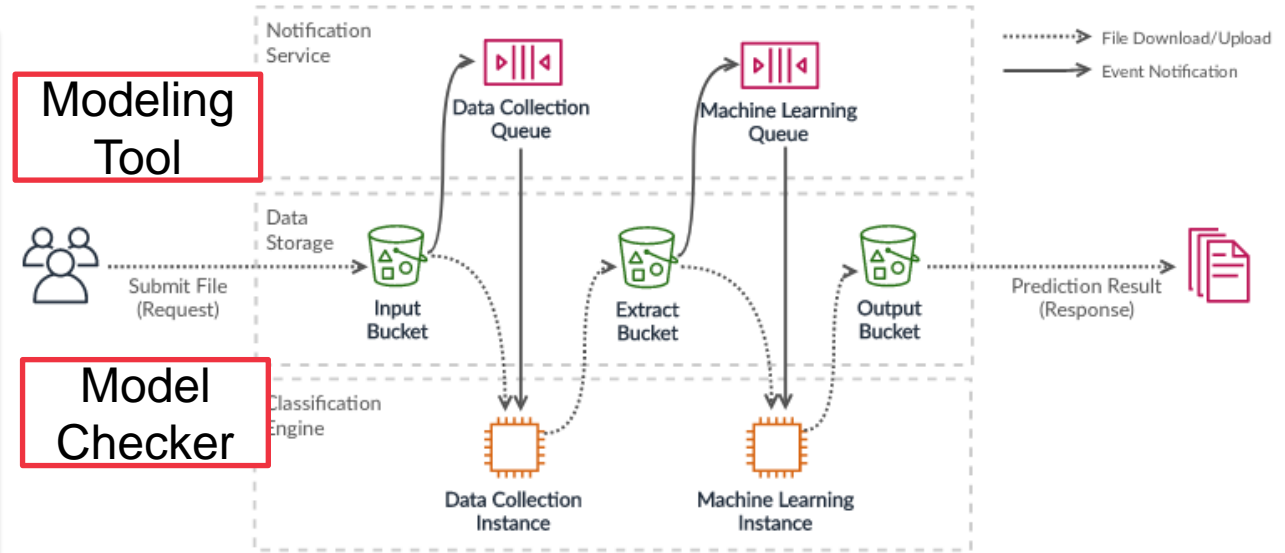
- **Responsive:** quality of services
- **Resilient:** deal within failures
- **Elastic:** deal with different workload and quality of analytics
- **Message-driven:** allow loosely coupling, isolation, asynchronous

Reactive systems for Big Data/ML: methods

- **Have different components as services**
 - components can come from different software stacks
 - components for doing computation as well as for data exchange
- **Elastic computing platforms**
 - platforms should be deployed on-demand in an easy way
- **Using messages to trigger tasks carried out by services**
 - messages for states and controls as well as for data
 - heavily relying on message brokers and lightweight triggers/controls (e.g., with serverless/function-as-a-service)

Examples: do-it-yourself ML classification for BIM

- Discussion: dealing R3E with ML workflows?
- Where, What, When and How



Source: Minjung Ryu, „Machine Learning-based Classification System for Building Information Models“, Aalto CS Master thesis, 2020

Which frameworks

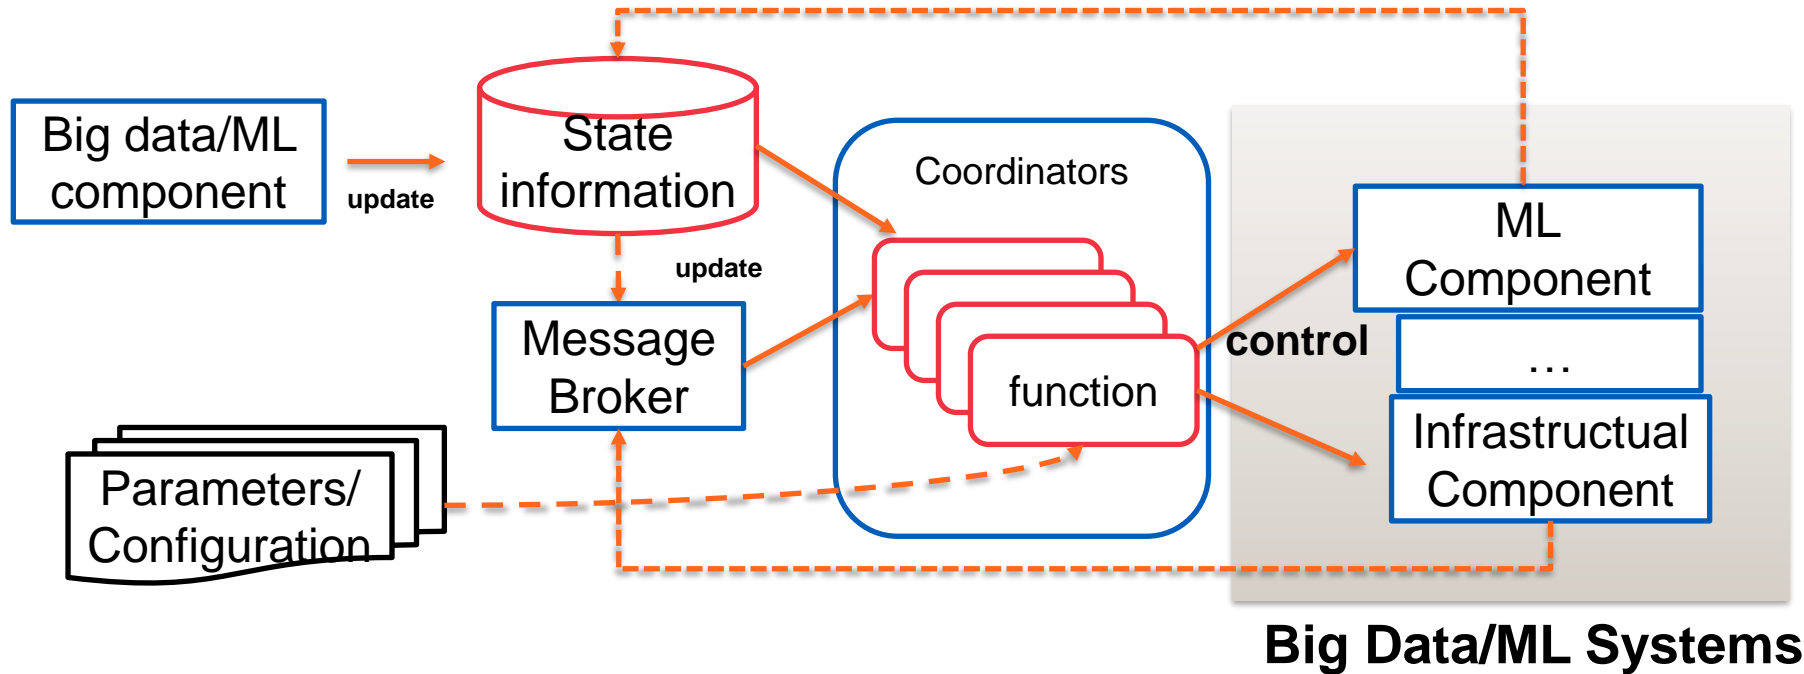
- **Message brokers**

- Kafka, RabbitMQ, Amazon SQS, ...
- types of messages and semantics must be defined clearly

- **Triggers and controls**

- serverless/function-as-a-service can be used: trigger a function based on a message
 - *AWS Lambda, Google Cloud Function, Knative, Kubeless, OpenFaaS, Azure Functions*
 - *We are discussing to use serverless for “coordination”*
- self-implementation of triggers listening messages

Common architecture



Example: Serverless for coordination

- **Training preparation**
 - before running a training: you move data from sources to stage, ship the code and prepare the environment
- **Coordination of ML phases**
 - do the coordination of three phases: data preprocessing, training and take the best model to deploy to a serving platform
- **Experiment results gathering**
 - you run experiments in different places. There are several logs of results, you gather them and put the result into a database

Serverless as functions within ML workflows

- Tasks in ML can be implemented as a function
- Thus a workflow of functions can be used to implement ML pipelines
 - Example: using serverless to implement data preprocessing

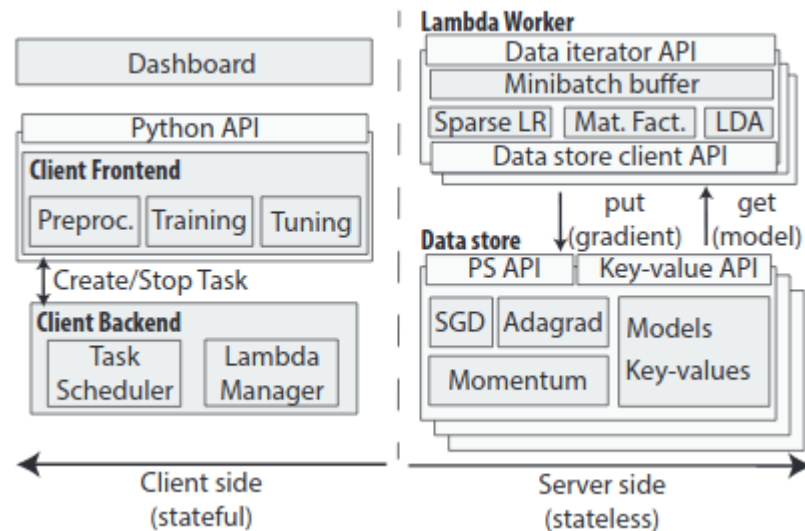


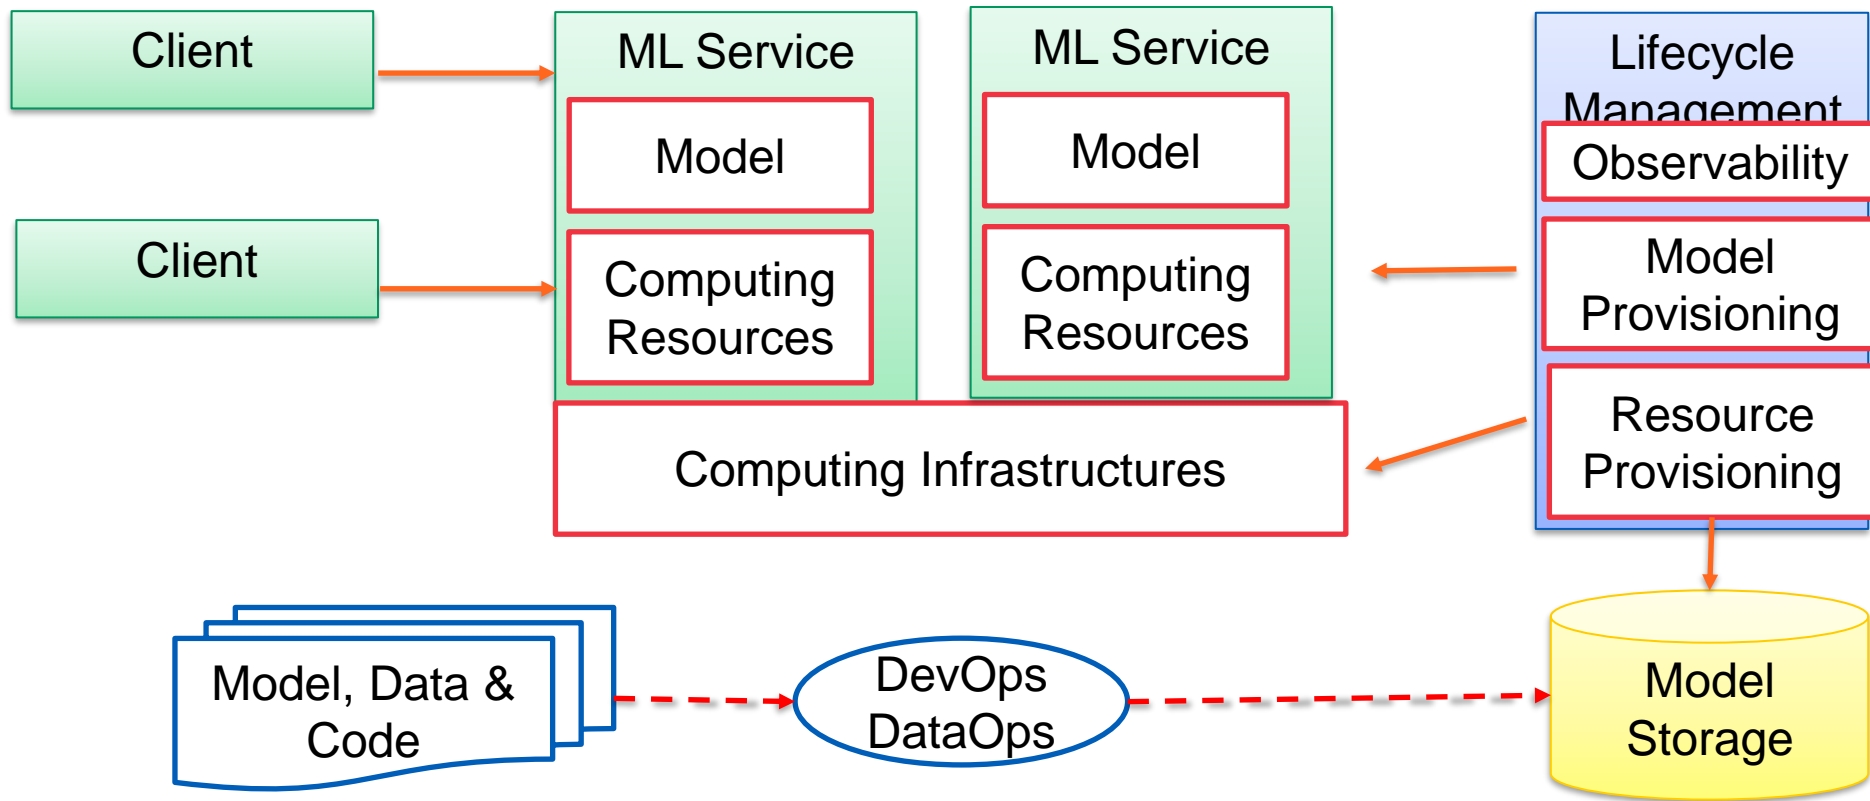
Figure source: Joao Carreira, Pedro Fonseca, Alexey Tumanov, Andrew Zhang, and Randy Katz. 2019. Cirrus: a Serverless Framework for End-to-end ML Workflows. In Proceedings of the ACM Symposium on Cloud Computing (SoCC '19). DOI:<https://doi.org/10.1145/3357223.3362711>

Dynamic ML Serving

ML Model Serving

- **Allow different versions of ML models to be provisioned**
 - runtime deployment/provisioning of models
 - “model as code” or “model as a service” → can be deployed into a hosting environment
- **Why? Anything related to R3E?**
 - concurrent deployments with different SLAs
 - A/B testing and continuous delivery for ML (<https://martinfowler.com/articles/cd4ml.html>)
- **Existing platforms**
 - increasingly support by different vendors as a concept of “AI as a service” (check <https://github.com/EthicalML/awesome-production-machine-learning#model-deployment-and-orchestration-frameworks>)

ML Model Serving design



ML Service

- **Long runtime inferencing services**
 - with well defined interfaces, know how to invoke models
 - accept continuous requests and serve in near-real time
- **Containerized service with REST/gRPC**
 - for on-demand serving or for scaling long running serving
- **Serverless function wrapping models**
 - short serving time
- **Batch serving**
 - not near real time serving due to the long inferencing time

Question: which forms are the best for which situations

Example: TensorFlow Extended Serving

- **Lifecycle**
 - Load, serving and unloading
- **Metrics & Policies**

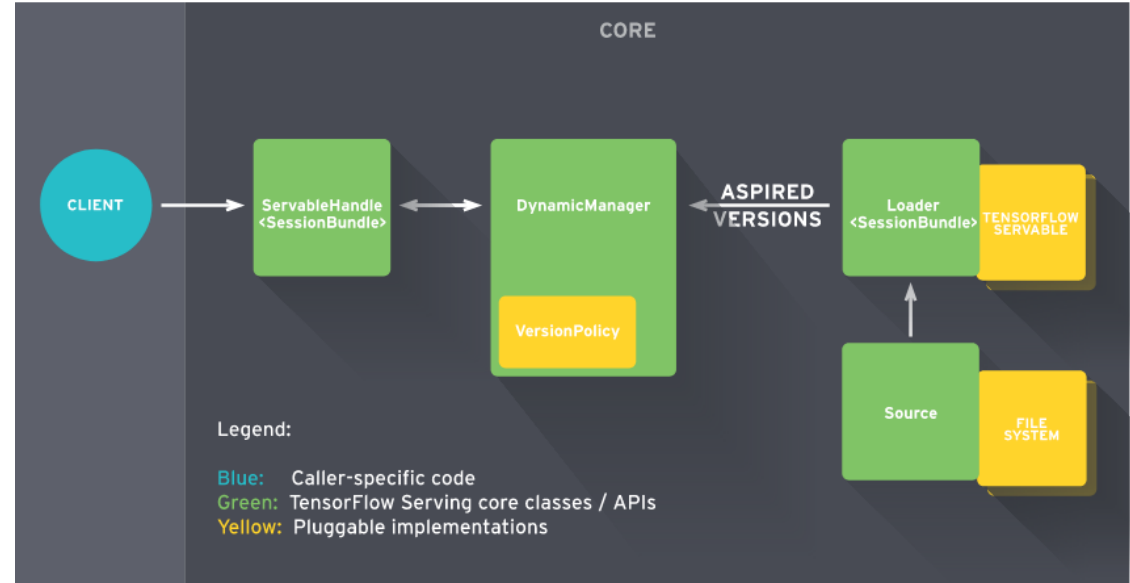


Figure source: <https://www.tensorflow.org/tfx/serving/architecture>

Example of Prediction.io

- Discussion: dealing R3E with ML workflows?
 - Elastic components?
 - Where, What, When and How

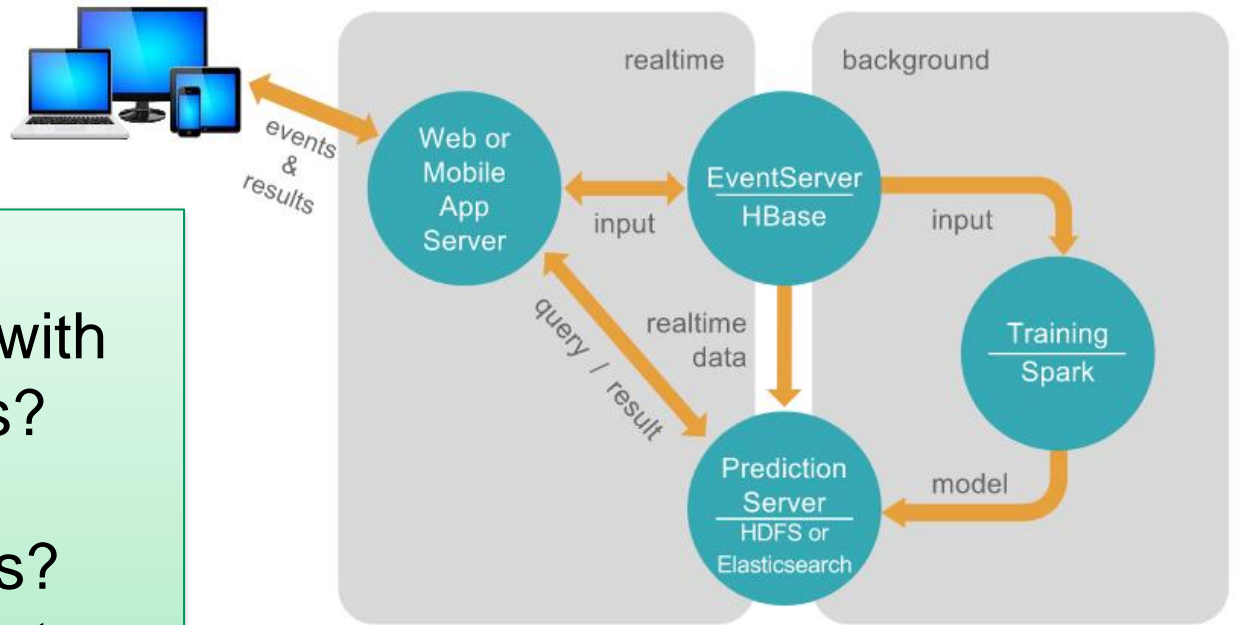
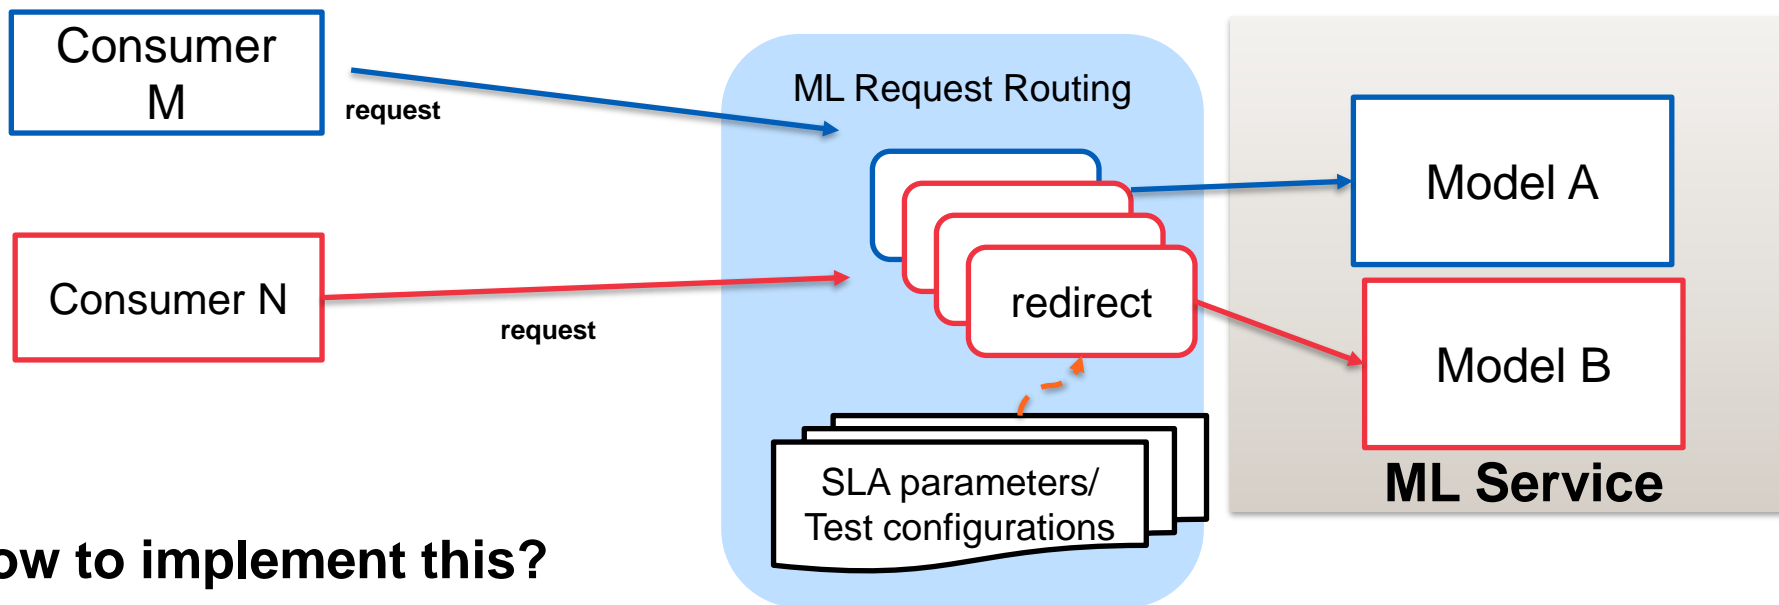


Figure source: <https://predictionio.apache.org/system/>

A/B testing and SLA-based serving

- Different models with different qualities/SLAs



How to implement this?

Example in Amazon Sagemaker

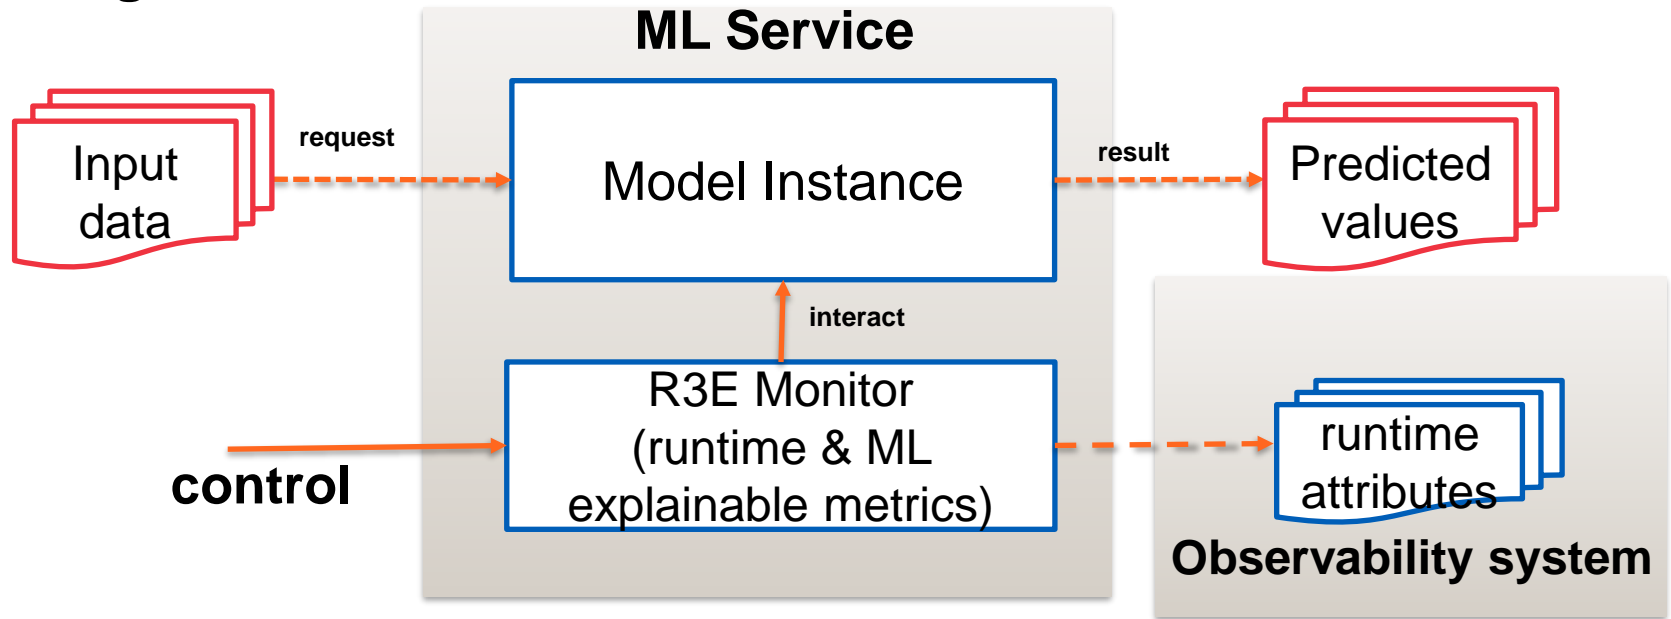
<https://docs.aws.amazon.com/sagemaker/latest/dg/model-ab-testing.html>

Load balancing/scaling model serving

- **ML inferencing capability in a ML model is encapsulated into a microservice or a task**
- **As a service**
 - with well-defined APIs (e.g., REST, gRPC), e.g., Dockerized service
 - using load balancing and orchestration techniques, such as Kubernetes
- **As a task**
 - using workflow management techniques to trigger new tasks
 - support scheduling, failure management and performance optimization by leveraging batch processing techniques

R3E Runtime attributes?

How to capture important metrics for observability and dynamic serving?



What if a model is too big, need a lot of computing resources?

Study log

P1 - Take one of the following aspects:

- P1.1 - Robustness, Reliability, Resilience or Elasticity
- P1.2 – Automation management

P2 - Check one of the following aspects:

- Orchestration or ML model serving

In a *specific software framework (F3)* that you find interesting/relevant to your work:

discuss how do you see F3 supports P1 in doing P2

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io