




Aalto University
School of Science

Coordination of Big Data/ML Tasks

Hong-Linh Truong
Department of Computer Science
linh.truong@aalto.fi, <https://rdsea.github.io>

Our focus in this course

The focus



**Any idea,
thought,
expectation?**

Content

- **Pipeline coordination**
 - Orchestration style
 - Choreography style
- **ML Model Serving**
- **Experiment Management**
- **Study log**

Pipeline coordination

Examples of Requirements

Discussion:

- **ML Phases & Tasks**
- **Software stack**
- **Execution environments**
 - Computing resources
- **R3E**

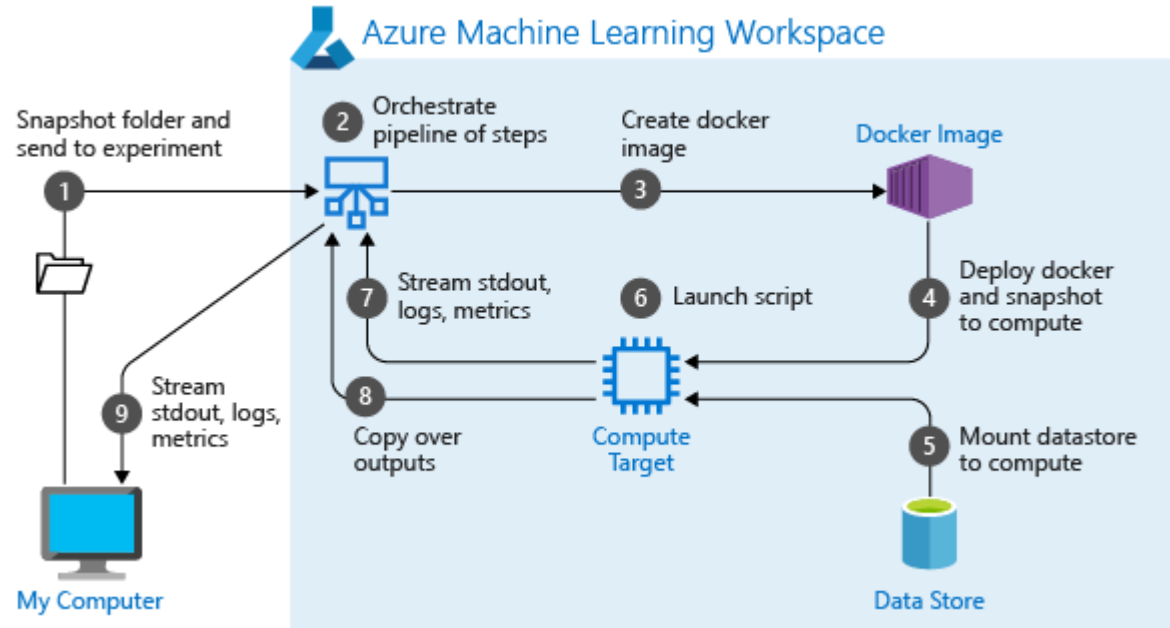
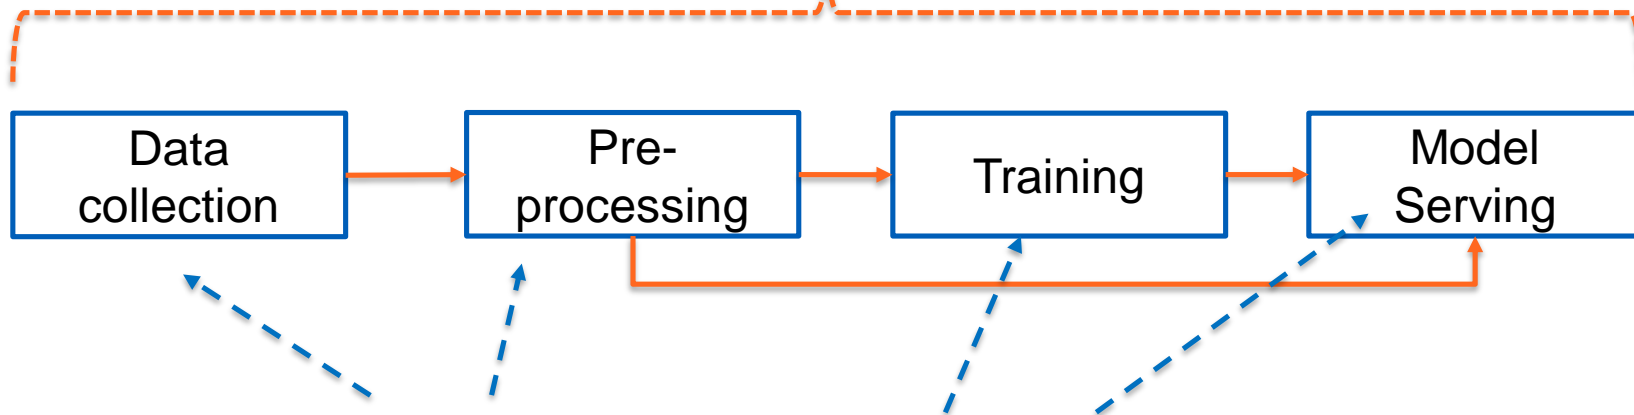


Figure source: <https://docs.microsoft.com/en-us/azure/machine-learning/concept-ml-pipelines>

The pipeline view of big data/ML systems

- **Multiple levels:**
 - Meta-workflow or -pipeline
 - Inside each phase: pipeline/workflow or other types of programs

(Meta) pipeline/workflow



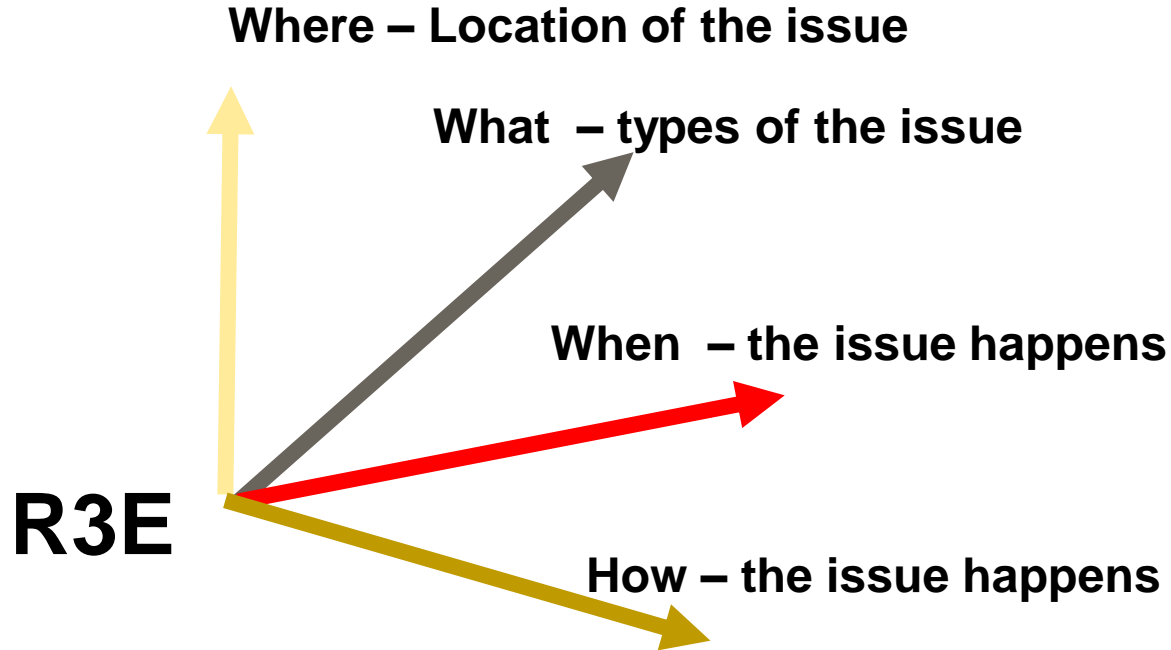
Airflow, function-a-as-service, Spark, Tensorflow, Keras, PyTorch,...



Main issues related to coordination

- **How to coordinate phases and tasks in big data/ML systems**
 - Automation is an important requirement
- **How to assure R3E for the pipeline execution**
 - End-to-end R3E requires coordination
 - Issues in internal and external services
- **How to manage experiment data**
 - Trial computing configurations
 - Inputs/results


W3H: what, when, where and how for R3E issues



Key notions

- **Workflow and Task/Activity/Step**
- **Important notes on the abstraction**
 - A task can encapsulate a “complex workflow”
- **Software frameworks**
- **Platform services**
 - Services offering features/functionality for executing “tasks”
 - Single or multiple the providers?
- **Execution environments and resources**
 - Single platform or cross (heterogeneous) platforms

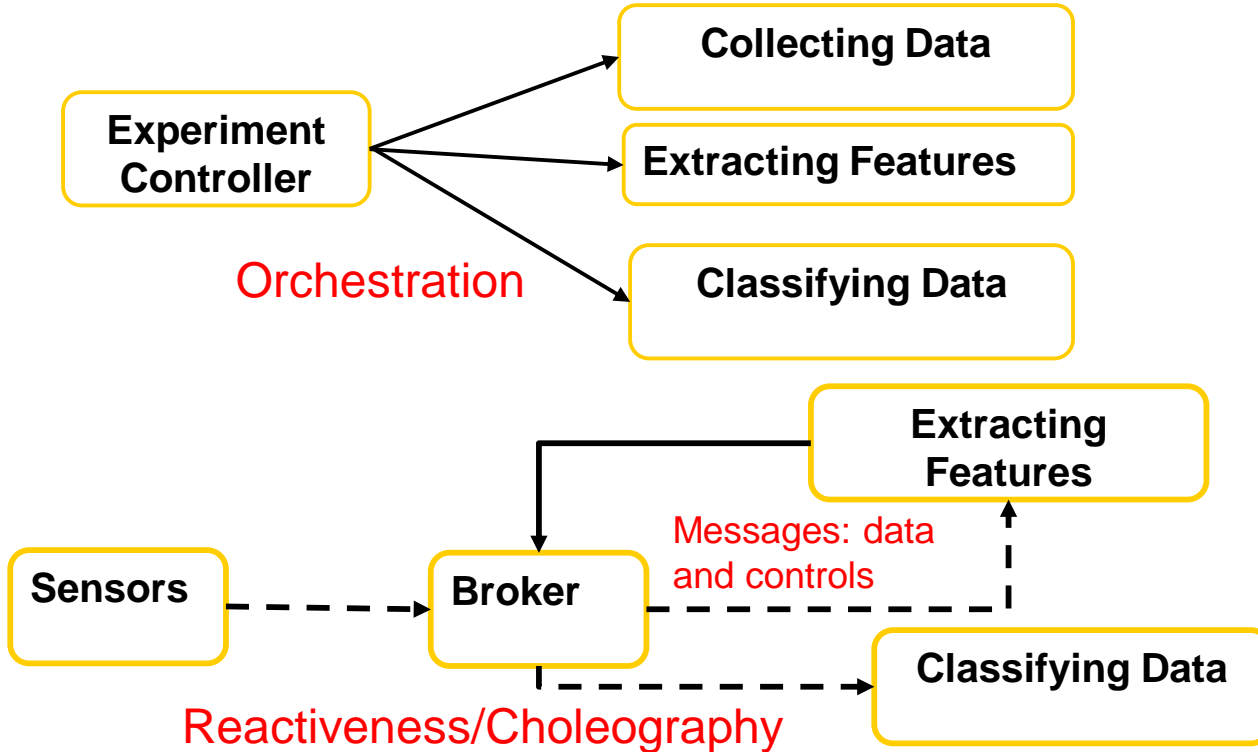
A task encapsulating a workflow via API



Coordination Styles

- **Coordination models for Big Data/ML systems**
 - Orchestration and reactivity/choreography
- **Orchestration**
 - Task graphs and dependencies are based on control or data
 - Triggered based on completeness of tasks or the availability of data
- **Reactivity/Choreography**
 - Follow reactive model: tasks are reacted/triggered based on messages

Orchestration and Reactiveness



System issues impacting R3E and coordination

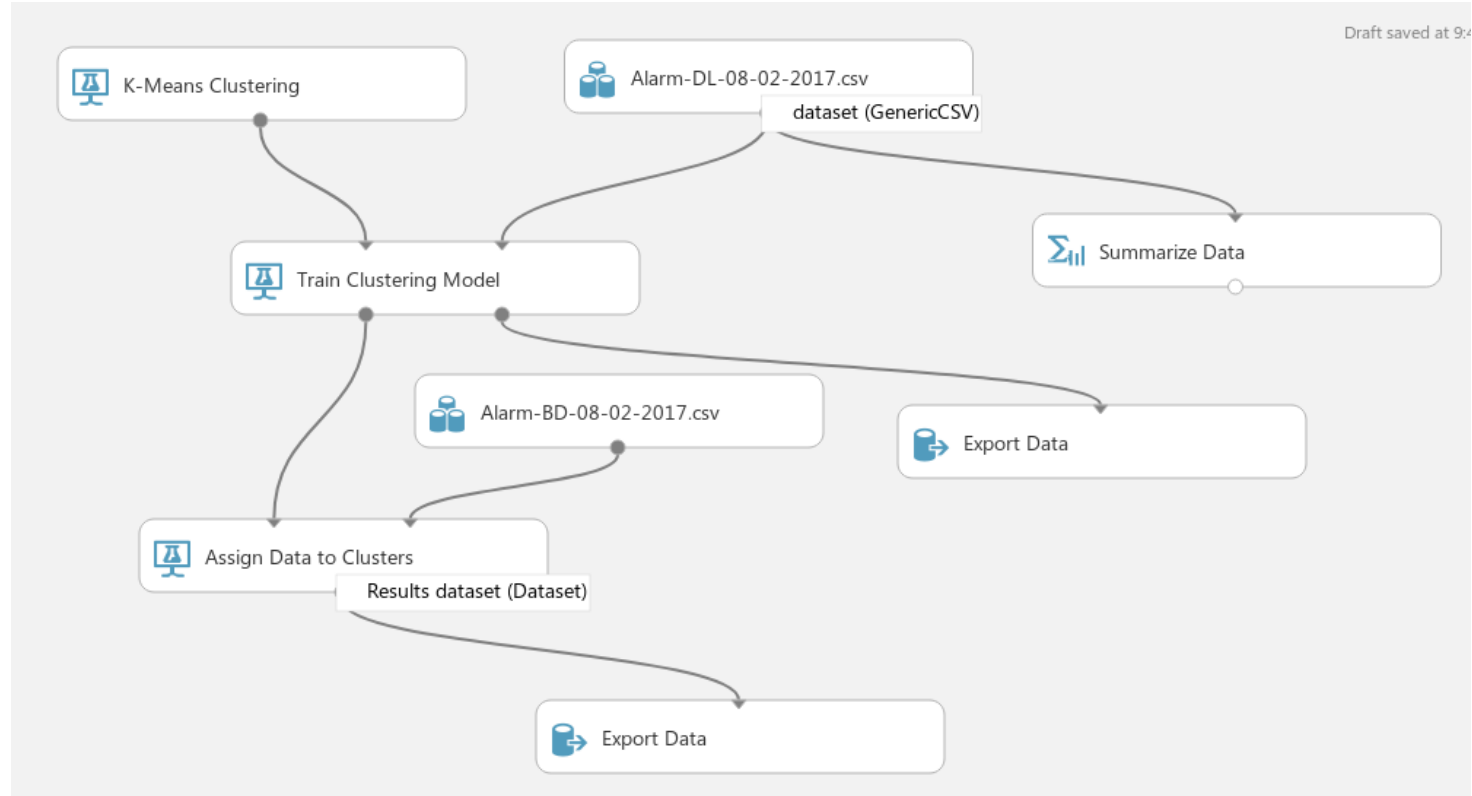
- **Main situations:**
 - Within the same system/infrastructure
 - *All services and computing resources belong to the same platform/infrastructure*
 - *E.g., running everything with Google Cloud or Microsoft Azure*
 - Across systems/infrastructures
 - *Services in different clouds or cloud data centers*
 - *E.g., Edge-cloud infrastructures*
 - The same software stack or not?
- **How such situations would affect the coordination/R3E?**

Workflows

- **Examples like**
 - Apache Airflow, Azure ML Pipeline
- **Often running in the same infrastructure**
- **Task-driven or data-driven specification**
- **Generic workflows**
 - Use to implement different tasks, such as machine provisioning, service calls, data retrieval
 - *Examples: Airflow, Argo Workflows*
- **Specific workflows for specific purposes**
 - E.g., Kubeflow (<https://github.com/kubeflow/pipelines>)

Workflow used in ML pipelines

Azure ML Pipelines




Orchestration architectural style: design

- **Workflow architectures are known**
 - Big Data/ML systems: leverage many types of services and cloud technologies
- **Required components**
 - Workflow/Pipeline specifications/languages (also UI)
 - Data and computing resource management
 - Orchestration engines (with different types of schedulers)
- **Execution environments**
 - Cloud platforms (e.g., VMs, containers, Kubernetes)
 - External services

Examples: Kubeflows

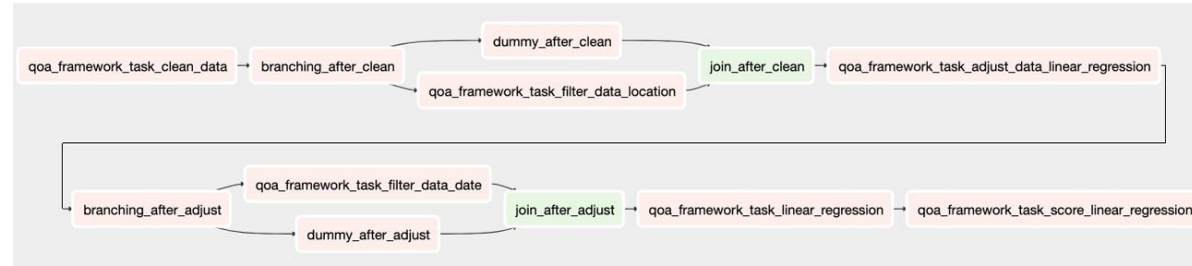
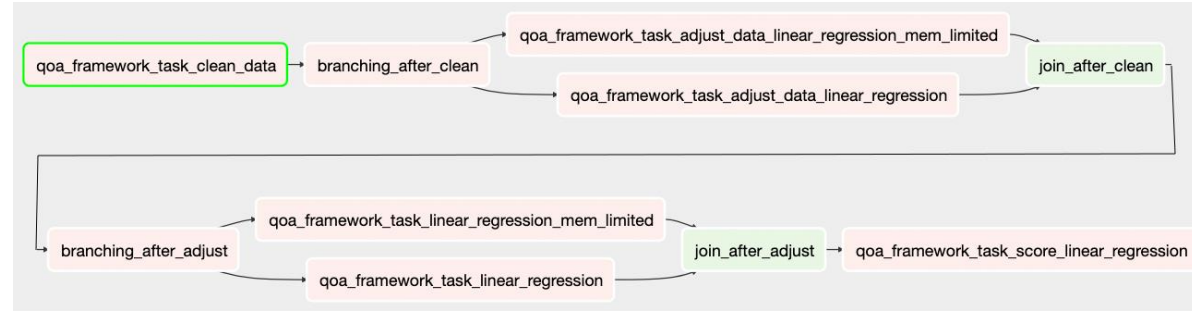
- End-to-end Orchestration
 - Orchestration is based on workflows
 - Using “Orchestration controllers”
- Discussion: dealing R3E with ML workflows?
 - Where, What, When and How



<https://www.kubeflow.org/docs/pipelines/overview/pipelines-overview/>

Examples: Coordinating tasks

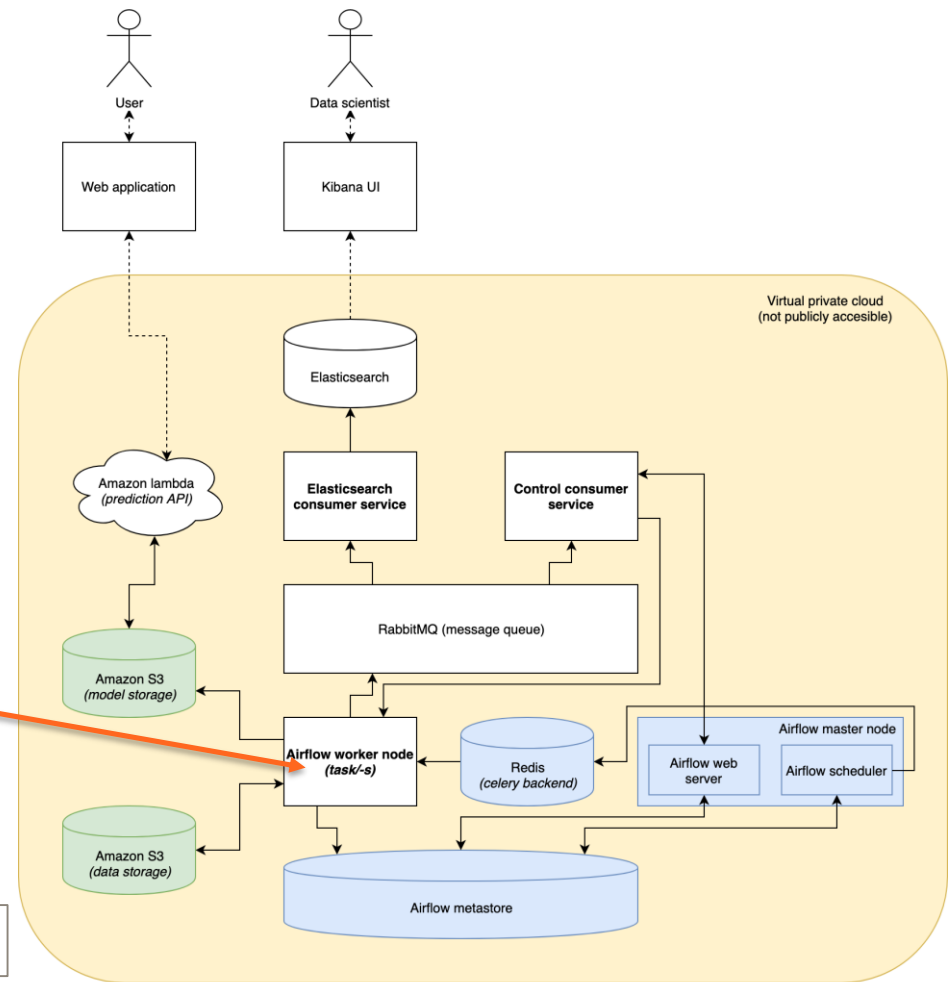
- Discussion: dealing R3E with ML workflows?
- Where, What, When and How



Source: Kreics Kristis, „Quality of analytics management of data pipelines for retail forecasting,“, Aalto CS Master thesis, 2019

Examples: Exchanging metrics for R3E coordination

Monitoring various metrics,
including user-defined
quality of data




Source: Kreics Kristis, „Quality of analytics management of data pipelines for retail forecasting“, Aalto CS Master thesis, 2019

Choreography: Reactive systems for Big Data/ML

Do you remember key principles of reactive systems?

Reactive systems



Source: <https://www.reactivemanifesto.org/>

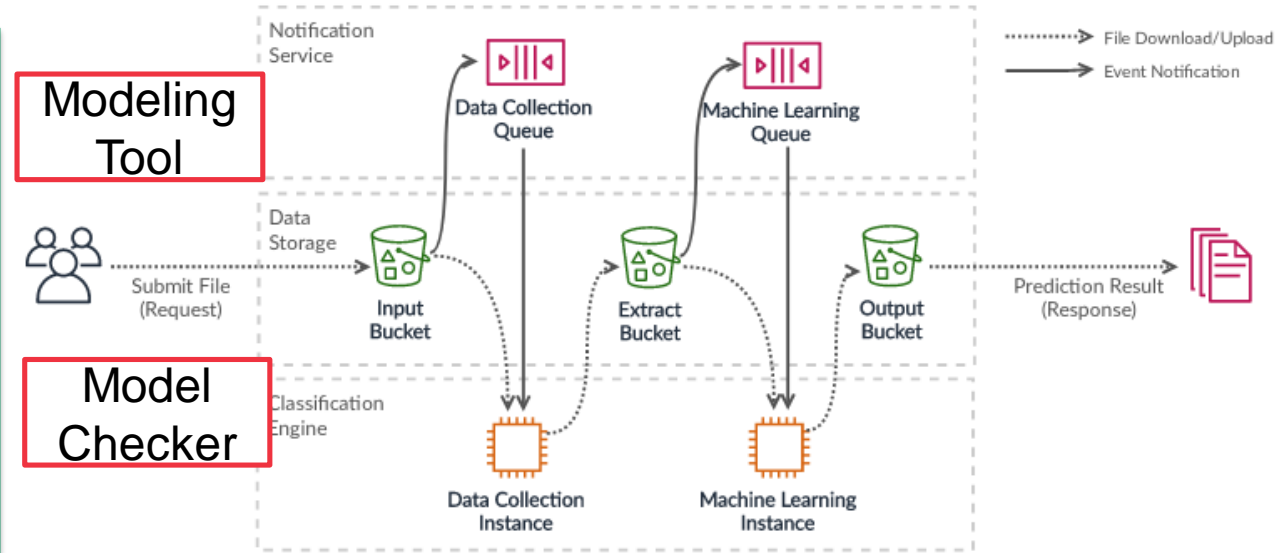
- **Responsive:** quality of services
- **Resilient:** deal within failures
- **Elastic:** deal with different workload and quality of analytics
- **Message-driven:** allow loosely coupling, isolation, asynchronous

Reactive systems for Big Data/ML: methods

- **Have different components as services**
 - Components can come from different software stacks
- **Elastic computing platforms**
 - Platforms should be deployed on-demand in an easy way
- **Using messages to trigger tasks carried out by services**
 - Messages for controls and for data
 - Heavily relying on message brokers and serverless (function-as-a-service)

Examples: do-it-yourself ML classification for BIM

- Discussion: dealing R3E with ML workflows?
- Where, What, When and How



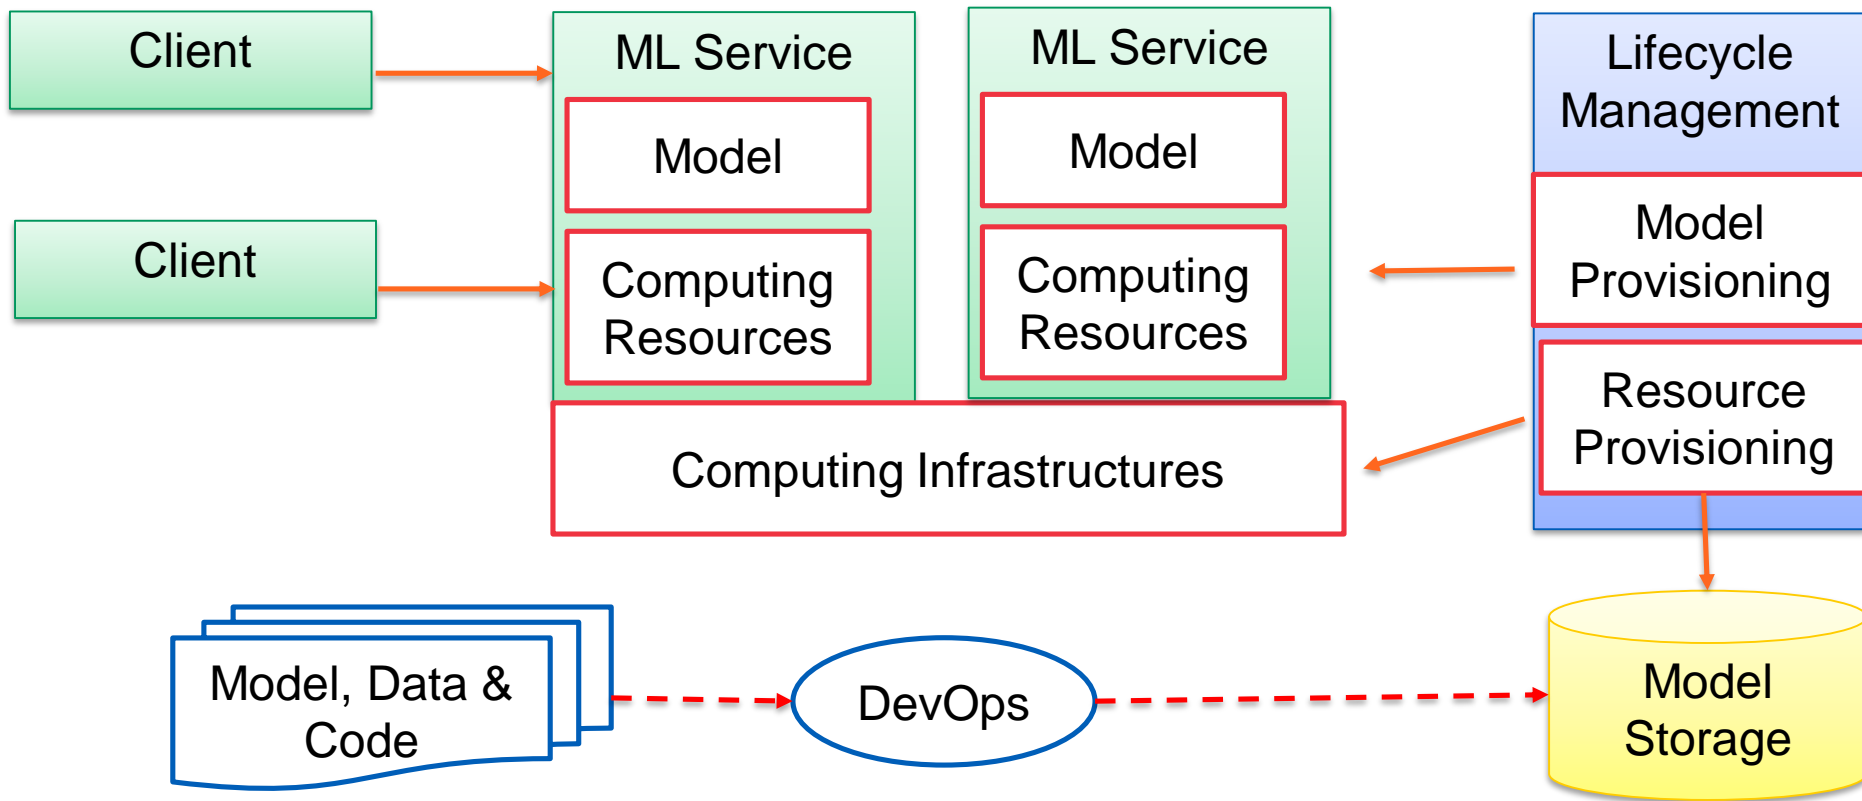
Source: Minjung Ryu, „Machine Learning-based Classification System for Building Information Models“, Aalto CS Master thesis, 2020

Dynamic ML Serving

ML Model Serving

- **Allow different versions of ML models to be provisioned**
 - Runtime deployment/provisioning of models
 - “Model as code” → can be deployed into a hosting environment
- **Why? Anything related to R3E?**
 - Concurrent deployments with different “pay-per-use” principles (elasticity as well?)
 - A/B testing and continuous delivery for ML (<https://martinfowler.com/articles/cd4ml.html>)
- **Existing platforms**
 - Increasingly support by different vendors as a concept of “AI as a service” (check <https://github.com/EthicalML/awesome-production-machine-learning#model-deployment-and-orchestration-frameworks>)

ML Model Serving design



Example: TensorFlow Extended Serving

- **Lifecycle**
 - Load, serving and unloading
- **Metrics & Policies**

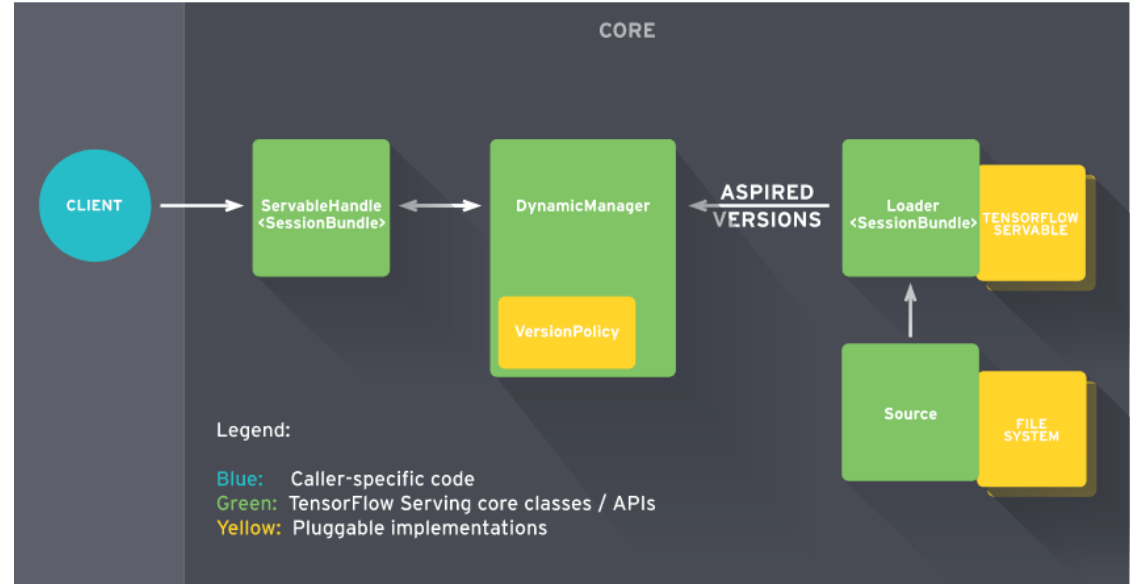


Figure source: <https://www.tensorflow.org/tfx/serving/architecture>

Example of Prediction.io

- Discussion:
dealing R3E with
ML workflows?
 - Elastic objects?
 - Where, What,
When and How

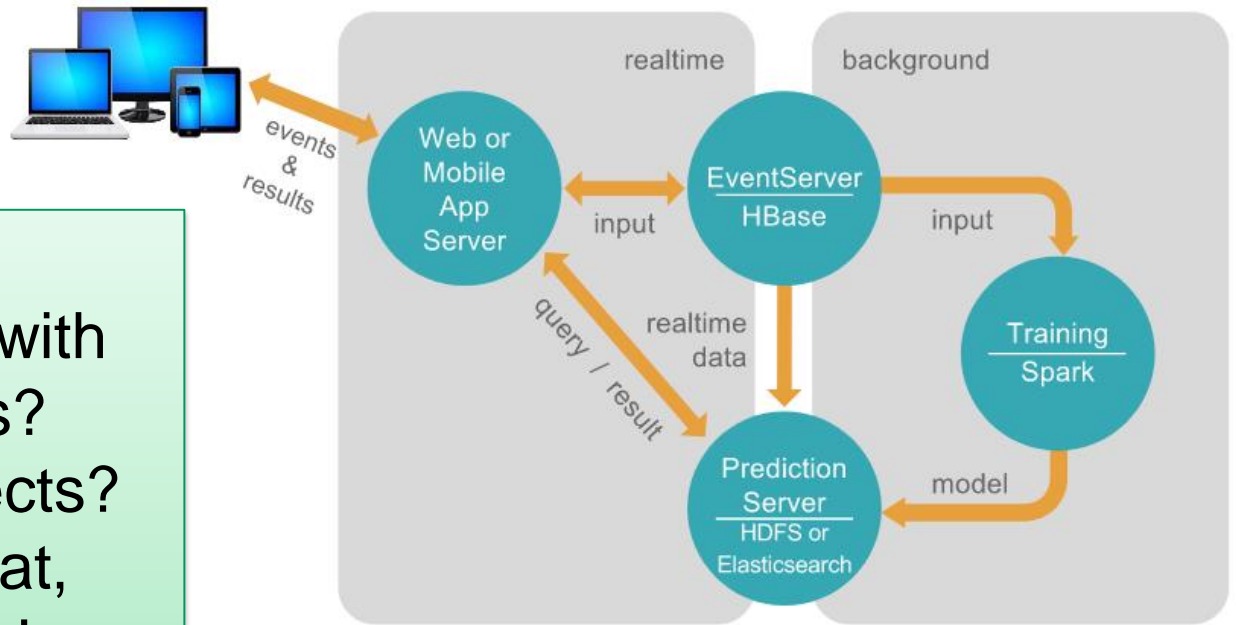


Figure source: <https://predictionio.apache.org/system/>

Experiment management

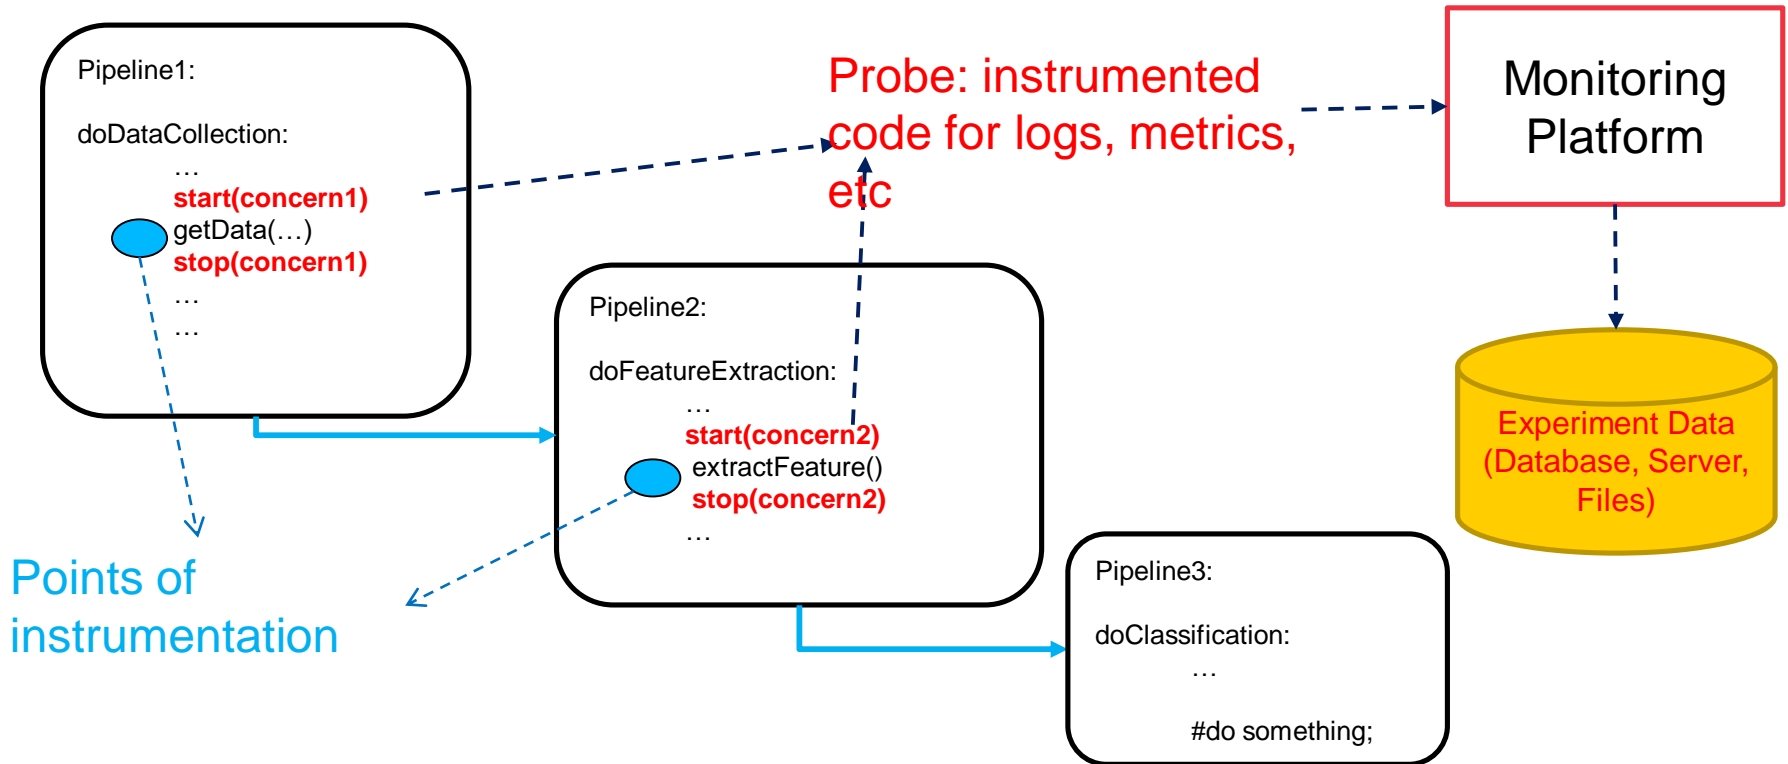
Problems

- **We need to run many experiments**
 - Known domain and well-known books (e.g., “Design and Analysis of Experiments” by Douglas C. Montgomery)
 - How does it work in ML?
- **Coordinating runs in experiments is also an important task**
 - How does this help to understand and support R3E?
- **What do we need?**
 - Tools/frameworks for tracking experiments

Notions

- **A single run/trial**
 - Inputs, results, required software artefacts
 - Computing resources, logs/metrics
- **Experiment**
 - A collection of runs/trials/executions gathered in a [specific context](#)
- **Steps**
 - Parameterization: generate different parameters
 - Deployment: prepare suitable environments
 - Execution: run and collect metrics
 - Analysis and Sharing: analyze experiment data

Experiment tracking



But remember it is very large system! Which tools can we use?

Examples

- **Experiment in Azure ML SDK**
 - <https://docs.microsoft.com/en-us/python/api/overview/azure/ml/?view=azure-ml-py#experiment>
- **MLFlows**
 - <https://mlflow.org/>
- **Kubeflows**
 - <https://www.kubeflow.org/docs/pipelines/overview/concepts/>

Examples: MLFlow APIs

- **Experiment**

```
mflow.start_run() / end_run()
```

- **Logs/metrics collection**

```
mflow.set_tag()
```

```
mflow.log_*()
```

- **Tracking data management**

- Local files, Databases, HTTP server, Databrick logs

Study log

P1 - Take one of the following aspects:

- Robustness, Reliability, Resilience or Elasticity

P2 - Check one of the following aspects:

- Orchestration, ML model serving or Experiment Management

In a *specific software framework* (F3) that you find interesting/relevant to your work:

discuss how do you see F3 supports P1 in doing P2

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io